



## Design Concepts

By: Ahmad



1



## Intelligent Traffic Management

[www.cognixia.com](http://www.cognixia.com)

2

# Intelligent Traffic Management (ITM)

- Intelligent Traffic Management **refers** to the **automated distribution** and **routing** of **network traffic** (requests, users, or data) **across multiple servers, services, or locations** to **optimize performance, availability, and reliability**.
- It's "intelligent" because it **can make decisions dynamically** based on **real-time metrics** like **server load, latency, location, or availability**.



[www.cognixia.com](http://www.cognixia.com)

3

## What is Load Balancer

- Load balancing is a mechanism to **distribute incoming network traffic** across **multiple servers** or resources to **prevent overload** on **any single server** and ensure **high availability** and **reliability**.
- Key Functions**
  - Distribute traffic
  - Monitor server health
  - Provide fault tolerance
  - Improve performance and user experience
- Types of Load Balancers**
  - Hardware Load Balancer**
    - Dedicated appliance (e.g., F5, Citrix NetScaler).
  - Software Load Balancer**
    - Runs as software (e.g., NGINX, HAProxy).
  - Cloud Load Balancer**
    - Managed service (e.g., AWS ELB, Azure Load Balancer).



[www.cognixia.com](http://www.cognixia.com)

4


## Proxy Mode and Pass-Through Mode

[www.cognixia.com](http://www.cognixia.com)

5

## Proxy vs. Pass-Through Modes

- The **operational mode** dictates how the **load balancer handles the client connection and network packets**.
- This is fundamentally **tied to the OSI layer** the LB operates on: **Layer 4 (L4)** or **Layer 7 (L7)**.
- Types**
  - Proxy Load Balancer Modes
  - Pass-Through Load Balancer Modes
- The **choice between Proxy Mode** (*typically Layer 7*) and **Pass-Through Mode** (*typically Layer 4*) in load balancing depends entirely on the **features required by your application** and the performance demanded by your traffic.



[www.cognixia.com](http://www.cognixia.com)

6

## Load Balancer Strategies

[www.cognixia.com](http://www.cognixia.com)

7

## Load Balancing Strategies

- Load balancing **strategies** can be broadly **categorized** as **Static** (rule-based) or **Dynamic** (real-time and adaptive).
- ITM solutions typically **rely** heavily on **dynamic methods**, often **incorporating** Artificial Intelligence (**AI**) and Machine Learning (**ML**).



[www.cognixia.com](http://www.cognixia.com)

8

# Static Load Balancing Algorithms



- **Static algorithms** use **pre-set rules** and do not account for the current health or load of a server, though some can be "weighted" to account for capacity differences.
  - **Round Robin**
    - Requests are **distributed sequentially** to **each server** in the pool.
    - It's simple but fails to account for varying server capacities or current workloads.
  - **Weighted Round Robin (WRR)**
    - Servers are **assigned weights** based on their processing power or capacity.
    - Servers with **higher weights** receive a proportionally **greater number of requests**.
  - **IP Hash (Source IP Hash)**
    - The **load balancer** performs a **calculation** (hash) on the **client's IP address**.
    - This result **consistently maps** that **client's requests** to the **same server**.
    - This is vital for **session persistence** where a user must **remain connected** to the **same server**.

[www.cognixia.com](http://www.cognixia.com)

9

# Dynamic Load Balancing Algorithms



- Dynamic algorithms use **real-time monitoring** of **server health** and **performance metrics** to make routing decisions on the fly, making them foundational to Intelligent Traffic Management.
  - **Least Connection**
    - Directs **new requests** to the **server** that **currently** has the **fewest active connections**.
    - This is highly effective for distributing uneven workloads.
  - **Weighted Least Connection**
    - An extension of **Least Connection** where the **server** with the **lowest ratio of active connections** to its **assigned capacity weight** receives the new request.
  - **Least Response Time (or Least Time)**
    - Routes **traffic** to the **server** that is **currently showing the fastest response time** (often measured by combining active connections and response time to a health check).
    - This is a strong strategy for optimizing user experience and minimizing latency.
  - **Resource-Based (Adaptive)**
    - This advanced method uses **specialized software** agents on **each server** to report **real-time metrics** like **CPU usage, memory consumption, or available bandwidth**.
    - The load balancer then **routes traffic** to the **server** with the most **available free resources**.

[www.cognixia.com](http://www.cognixia.com)

10



## Improving Performance

[www.cognixia.com](http://www.cognixia.com)

11

## What is Caching

- **Caching** is the process of **storing frequently accessed data in a faster, temporary storage location (cache)** so that future requests for the same data can be served faster—**without repeatedly fetching it from the original slow source** (database, disk, API, backend service).
- **Why Use Caching**
  - **Faster Response:** Reduces time to serve data (*microseconds instead of seconds*)
  - **Reduces Load on Backend:** Cuts down database / API calls
  - **Cost Optimization:** Fewer compute and **database costs**
  - **Scalability:** Handles more traffic without increasing resources
  - **Better User Experience:** Faster loading pages, smoother apps




[www.cognixia.com](http://www.cognixia.com)

12

## Caching ... continue

- **How Caching Works**
  - Client Request → Cache? →
    - ✓ Hit → Return Cached Data (Fast)
    - ✗ Miss → Fetch from DB/API → Store in Cache → Return to Client
- **Service Providers**
  - **CDN**: Akamai, Cloudflare, Fastly, Edgio, StackPath, etc.
  - **In-Memory**: Redis, Memcached, Hazelcast, etc.



[www.cognixia.com](http://www.cognixia.com)

13

## Caching Types

[www.cognixia.com](http://www.cognixia.com)

14


## Cache | Types ... continue



- Caches are **classified based** on where they **physically reside** within the system architecture:

- o CDN Caching (Content Delivery Network)**

- Data stored on geographically distributed servers (PoPs).**
    - It reduces the physical distance** content must travel, serving high-volume, global traffic.
    - When a **user requests** an **asset**, the **request is routed** to the **nearest PoP**.
    - If the **asset** is in the **PoP's cache (cache hit)**, it's **served immediately**.
    - If not (**cache miss**), the **PoP fetches** it from the **original web server** (the origin) and **caches** it for **future requests**.



www.cognixia.com


15

## Cache | Types



- Distributed Cache (Cluster cache)**

- A **separate, dedicated cluster** of high-speed **servers** (e.g., Redis, Memcached) **accessible** over the **network**.
  - In-Memory Cache**
    - The term In-Memory Cache **refers** to a **system** where **data is stored** entirely in **RAM** of **dedicated servers**.
    - Data stored in RAM** is **accessed** orders of magnitude **faster than data stored on disks** (even SSDs).
  - Dynamic application data such as **user session data**, **database query results**, and computed **API responses**.
  - Unlike databases that store complex rows and tables, in-memory **caches** usually **store data** in **simple key-value pairs**, allowing for extremely **fast lookups** by key.



www.cognixia.com

16



## Caching Strategies

[www.cognixia.com](http://www.cognixia.com)

17

## Caching Strategies



- The strategy you choose **dictates** when **data** is **written** to the **cache** and how it is **updated**.
  - **Cache-Aside (Lazy Loading)**
  - **Write-Through**


[www.cognixia.com](http://www.cognixia.com)

18

# Caching Strategies



- The strategy you choose **dictates** when **data** is **written** to the **cache** and how it is **updated**.
- Cache-Aside (Lazy Loading)**
  - Mechanism:** The application or client is responsible for checking the cache first.
    - Cache Hit**
      - If the **data** is found **returned immediately**.
    - Cache Miss**
      - If the **data** is **not** in the **cache**, the **application fetches** it from the **database**, **returns** the data to the **client**, and then **writes** a **copy** of the data into the **cache** for future requests.
  - Pros**
    - Only requested data is cached**, preventing the cache from being filled with unused items.
  - Cons**
    - Higher latency** on the **first request** (*the cache miss*).
    - Data can become stale** until a subsequent request updates it.
  - Best For**
    - Read-heavy workloads** where data access patterns are unpredictable.



www.cognixia.com

19

# Caching Strategies ... continue



- Write-Through**
  - Mechanism**
    - Data is simultaneously written to both the cache and the primary database before the operation is considered complete.**
  - Pros**
    - Ensures data consistency** between the **cache** and the **database** at all times.
    - Reduces latency** for future reads.
  - Cons**
    - Higher latency** on write operations, as the **application** must wait for both the **cache** and the **database writes to complete**.
  - Best For**
    - Applications where data consistency** and **immediate read availability after a write** are critical.

## 3. Write-Through

Ensures data consistency on write operations.



www.cognixia.com

20




# Service Orchestration

[www.cognixia.com](http://www.cognixia.com)

21

## What is Service Orchestration

- Service Orchestration is a **method of arranging and coordinating multiple independent services** (often *microservices*) to **execute** a complex **business process or workflow**.
- It **dictates the flow of control and the sequence of actions** that must **occur across** various **services** to achieve a final goal.
- In essence, an **orchestrator** (*a dedicated service or component*) takes the **lead** and **tells each participant service exactly what to do and when to do it**.



[www.cognixia.com](http://www.cognixia.com)

22



## API Gateway

[www.cognixia.com](http://www.cognixia.com)

23

## API Gateway

- An **API Gateway** is a **central component** in modern software architectures (*especially microservices*) that **acts as a single point of entry** for all **external client requests**.
- It functions as a **reverse proxy** that **routes requests** to the appropriate **backend services**, while **simultaneously handling** common cross-cutting concerns like **security, caching, and traffic management**.
- **API Gateway** is the **single entry point** for all **client requests** into **backend services** (*web services, microservices, databases, authentication, etc.*) in modern architectures.



[www.cognixia.com](http://www.cognixia.com)

24

## API Gateway ... continue



- It acts as:
  - Reverse Proxy
  - Traffic Router
  - Security Gate
  - Orchestrator of service calls
- **Supported By:** Kong Gateway, Tyk API Gateway, Apigee Gateway, etc.
- **Purpose:**
  - The primary purpose of an API Gateway is to **decouple clients** (like web browsers or mobile apps) from the complex, evolving architecture of the backend services.

Purpose	Detailed Usage
Request Routing	Directs incoming requests to the correct internal microservice based on the URL path, host, or other criteria.
Security Enforcement	Handles centralized tasks like authentication (verifying API keys or tokens) and authorization, shielding backend services from external threats.
Traffic Management	Implements rate limiting, throttling, and load balancing across internal service instances.
Request Aggregation	Combines responses from multiple backend services into a single, cohesive response for the client (known as <b>API Composition</b> or <b>Backend For Frontend - BFF</b> ).
Quality of Service	Handles service discovery, logging, monitoring, and applying Quality of Service (QoS) controls.

25


## Service Mesh


[www.cognixia.com](http://www.cognixia.com)

26

# What is Service Mesh

- A Service Mesh is an **infrastructure layer** that manages **service-to-service communication** in a **microservices architecture**.
- Instead of **embedding networking logic** (*like retries, security, observability*) inside each service, the **mesh handles it transparently via sidecar proxies** (e.g., *Envoy*).
- Think of it as the “**traffic control system**” inside your cluster, ensuring **services talk to each other** reliably, securely, and observably.
- **Feature**
  - **Secure communication:** Encrypts service-to-service traffic using mTLS (Mutual TLS)
  - **Traffic management:** Load balancing, retries, timeouts, circuit breakers
  - **Observability:** Tracks metrics, logs, traces of each request
  - **Policy management:** Controls access rules, rate-limiting, quotas
  - **Resilience:** Ensures services keep running even during failures



27

## Service Mesh ... continue

- The **Service Mesh** architecture **consists of two main parts**:
- **Data Plane**
  - The actual network of **proxies that intercepts** and **handles all service traffic**.
  - **Mechanism**
    - ✓ Every instance of a microservice is paired with a **Sidecar Proxy**.
    - ✓ All **incoming and outgoing network traffic** for that service instance **goes through** its dedicated **sidecar proxy**.
  - **Examples:** Envoy (the most popular open-source sidecar proxy).
- **Control Plane**
  - The **central brain** that **manages, configures**, and **monitors** all the **sidecar proxies** in the **Data Plane**.
  - **Mechanism**
    - ✓ It converts **high-level policies** (*defined by the operator*) into concrete **configurations** (*like routing rules and mTLS certificates*) and **pushes them** out to the **sidecars**.
  - **Examples:** Istio, Linkerd, Consul Connect.

www.cognixia.com

28



# Session Management

[www.cognixia.com](http://www.cognixia.com)

29

## What is Session Management

- User Session Management is the **process** of **tracking a user's activity and status** (*their "state"*) across multiple, non-contiguous requests to a server.
- Since **HTTP is inherently stateless** (*each request is treated as brand new*), the **system must have a strategy to preserve the user's context** (e.g., *login status, shopping cart, preferences, role, authentication, etc.*).
- The **two primary architectural patterns** for doing this are **Stateful** and **Stateless** sessions.



[www.cognixia.com](http://www.cognixia.com)

30



## Stateful Session

[www.cognixia.com](http://www.cognixia.com)

31

## Stateful Session Management



- In a **Stateful system**, the **server** (or an *external store* managed by the server) is **responsible** for **retaining all information** about the **client session**.
- **How it Works**
  1. **Login**
    - The **user logs** in to the server.
  2. **Server Creates State**
    - The **server generates** a **unique**, unguessable **Session ID** (e.g., a random string like *aB7c4D9e*).
    - It **stores** all **session data** (*User ID, Role, Cart items*) locally in its **own memory**, a **database**, or a **dedicated session store** (like *Redis*).
  3. **Client Stores ID**
    - The **server sends** this **Session ID** back to the **client**, usually **embedded** in a **Cookie**.
  4. **Subsequent Request**
    - The **client sends** the **Cookie** (*containing the Session ID*) with **every subsequent request**.
  5. **Server Lookup**
    - The **server receives** the **ID** and must **perform a lookup** in its **session store** to **retrieve** the **user's full state**.

[www.cognixia.com](http://www.cognixia.com)

32



## Stateless Session

[www.cognixia.com](http://www.cognixia.com)

33

## Stateless Session Management



- In a **Stateless system**, the **client holds the state information**, which is **encoded** and **cryptographically signed** by the server. The **server retains no session record**.
- **How it Works (Using JWT)**
  - The most **common implementation** uses **JSON Web Tokens (JWT)**.
    1. **Login**
      - The **user logs** in to the server.
    2. **Server Creates Token**
      - The **server creates a JWT, embedding essential user claims** (e.g., `user_id`, `role`, `expiry_time`) **into the token's payload**.
      - The **server then signs the token** with a **secret key**.
    3. **Client Stores Token**
      - The **server sends** the complete, self-contained **JWT back to the client** (*usually in local storage or an HTTP header*).
    4. **Subsequent Request**
      - The **client sends** the **JWT with every subsequent request**.
    5. **Server Verification**
      - The **server receives** the **JWT** and simply **verifies** the **cryptographic signature** using its secret key.
      - If the **signature is valid**, the data inside the token is trusted, and the **server can process the request** without looking up any database.

[www.cognixia.com](http://www.cognixia.com)

34



## Fair Usage

[www.cognixia.com](http://www.cognixia.com)

35

## What is Fair Usage

- Fair usage ensures that **no single user, system, or service consumes more than their allowed share of resources**, preventing misuse, overloading, and ensuring availability, performance, and cost control for all users.
- This is commonly enforced using:
  - **Rate Limiting**
  - **Throttling**



[www.cognixia.com](http://www.cognixia.com)

36



# Rate Limiting

[www.cognixia.com](http://www.cognixia.com)

37

## Rate Limiting



- It **sets a strict limit** on the number of requests a user or system is allowed to make **within a time period** (per second, minute, hour, or day).
- If the limit is crossed, the system **blocks additional requests immediately**, usually returning **HTTP 429**
- **Focus:** Enforcement of a **hard quota or cap on usage**.
- **Goal:** **Protect the service from abuse**, such as **DDoS attacks**, brute-force login attempts, and unauthorized data scraping, by denying requests outright.
- **Purpose**
  - **Prevent Abuse** → **Stops API spamming**, bots, or DDoS attacks.
  - **Ensure Fair Usage** → **Divides API usage fairly** among all users.
  - **Control Costs** → **Protects backend resources** and cloud cost overruns.

[www.cognixia.com](http://www.cognixia.com)

38

19



# Throttling

[www.cognixia.com](http://www.cognixia.com)

39

## Throttling



- Throttling **slows down the request processing rate**, instead of blocking instantly.
- It **queues or delays** extra requests instead of rejecting them immediately.
- **Focus:** Regulation of the traffic flow to match the server's processing capacity.
- **Goal:** Manage traffic surges gracefully, maintain stable service performance during peak loads, and provide a better user experience by avoiding abrupt rejections.
- **Purpose:**
  - Avoid server overload → Smooths out request spikes instead of blocking users
  - Maintain performance → Ensures backend stays stable during peak load
  - Graceful Degradation → Allows requests but at slower speed

[www.cognixia.com](http://www.cognixia.com)

40


20

## How to Use Both Together



- Rate Limiting and Throttling can **both be used together** in a single program or system.
- **High-Level Logic:**
  1. Let the user send requests
  2. First apply **Rate Limiting**:
    - o If user exceeds allowed limit (e.g., more than 100 requests/minute) → **reject (429)**
  3. If user is within limit, but requests are coming too fast or system is overloaded → **Apply Throttling** → Delay or slow down processing

**Result:** Valid users are **slowed** (not blocked), while abusers are **completely blocked**.



[www.cognixia.com](http://www.cognixia.com)

41

## Asynchronous Communication

[www.cognixia.com](http://www.cognixia.com)

42

# What is Asynchronous Communication



- Asynchronous communication is a **system design approach** where a **sender sends a message or a request without waiting for an immediate response**.
- The **sender continues its work**, and the **receiver processes the message independently** at its own pace.
- Asynchronous communication means systems or components **do not wait for each other to respond**.
- They **send a message and continue their work**, allowing **high performance, loose coupling, scalability, and resilience**.
- Instead of calling services directly and waiting (synchronous), they communicate **via messages or events**.
- Analogy**
  - WhatsApp Message** → You **send a message and continue working**. The **receiver reads it later**.
  - Food Delivery App** → You **order food**, and **restaurant prepares it independently — you don't wait** on the phone.
  - Courier Service** → You **give a package to courier**; they **deliver it asynchronously**.


[www.cognixia.com](http://www.cognixia.com)

43

# Asynchronous Communication | Types



- The **types** are primarily **categorized** by the **communication model** used—how many receivers get the message and how the message is handled.
  - Message Queues**
    - (Point-to-Point)
  - Publish/Subscribe (Pub/Sub)**
    - (One-to-Many)
  - Event Streaming**
    - (Durable Log)
  - Webhook / Callback**
    - (HTTP Push)



[www.cognixia.com](http://www.cognixia.com)

44

# Message Queue

[www.cognixia.com](http://www.cognixia.com)

45

## Message Queues

- A **Message Queue (MQ)** is a form of asynchronous, point-to-point communication where a sender (**Producer**) places a message onto a **queue**, and a single recipient (**Consumer**) retrieves and processes it.
- **Key Characteristics**
  - **Communication Model**
    - **One-to-One (Point-to-Point).**
  - **Message Consumption**
    - A **message** is **consumed** by **exactly one** of the **consumers** **listening** to that **queue** and is then typically removed from the queue.
  - **Ordering**
    - Often maintains a **FIFO (First-In, First-Out)** order of messages for a single consumer.

### MESSAGE QUEUE: Point-to-Point Communication


 Task: Process Order

 One-to-One

 Task Distribution

[www.cognixia.com](http://www.cognixia.com)

46

## Message Queues ... continue



- **Technologies:** RabbitMQ, IBM MQ, IronMQ, Apache ActiveMQ, Beanstalkd, Gearman
- **Use Case**
  - **Order Processing in E-commerce:** Order service **puts orders** into a **queue**; **Inventory, Payment service reads it** one by one
  - **Resize Images in background:** User uploads image → **queued** → **worker resizes** asynchronously
  - **Email/SMS Notifications:** Message goes to queue and **notification service processes** it
- **Example**
  - **E-commerce Order Processing** (e.g., using **Rabbit MQ** or **IBM MQ**).
    - When a **user clicks "Place Order,"** the web service **immediately puts** an **OrderPlaced message** into a **queue** and **returns a confirmation** to the user.
    - A separate **fulfillment service picks up the message from the queue, processes** the **inventory check, payment, and shipping label generation**, without blocking the user interface.

[www.cognixia.com](http://www.cognixia.com)

47

## Event-Driven Architecture (EDA)



[www.cognixia.com](http://www.cognixia.com)

48

24

## What is Event-Driven Architecture (EDA)



- Event-Driven Architecture (EDA) is a **software design pattern** where the entire system's **behavior** is **centered** around the production, detection, consumption, and **reaction** to **events**.
- Instead of services making direct, synchronous calls to one **another** and **waiting** for a **response** (*the traditional request-response model*), they simply **announce** that a significant **change of state** (*an event*) has **occurred**.
- Other services that are interested in **that change** then **react asynchronously** and **independently**.
- The core goal of EDA is to achieve **loose coupling**, **high scalability**, and **real-time responsiveness**.
- **Core Models of EDA** (Communication Style)
  - **Publish/Subscribe (Pub/Sub)**
  - **Event Streaming**

[www.cognixia.com](http://www.cognixia.com)

49

## Publish/Subscribe (Pub/Sub)




[www.cognixia.com](http://www.cognixia.com)

50

# Publish/Subscribe (Pub/Sub)



- Publish/Subscribe (Pub/Sub) is a **broadcast model** where a sender (**Publisher**) sends a **message** (*called an event*) to a named **channel** or **Topic**, and all interested recipients (**Subscribers**) simultaneously receive a copy of that event.
- Unlike queues, **message is not consumed once**, it's delivered to **all subscribers**.
- **Key Characteristics**
  - **Communication Model**
    - **One-to-Many (Broadcast)**.
  - **Message Consumption**
    - A **message** is **delivered** to **all services subscribed** to that **topic**.
  - **Decoupling**
    - **Publishers and Subscribers have no knowledge of each other**, communicating only through the central broker/topic.



[www.cognixia.com](http://www.cognixia.com)

51




# Event Streaming

[www.cognixia.com](http://www.cognixia.com)

52

# Event Streaming

- **Event Streaming** is an advanced, powerful approach to data management that treats data as a **continuous, real-time, ordered, and durable stream of events**. It is one of the primary types of Event-Driven Architecture (EDA).
- **Event Streaming** is the continuous collection, processing, and delivery of real-time events (data) as they happen.
- Instead of storing and sending data later, systems stream events instantly to consumers that want to react immediately.
- It enables **Real-time processing, Live monitoring, Analytics, and Event-driven systems**.



[www.cognixia.com](http://www.cognixia.com)

53


## Webhook

[www.cognixia.com](http://www.cognixia.com)

54

## What is Webhook

- A **Webhook** is a **web-based callback**, where a system sends an **HTTP request (usually POST)** to your URL when a **specific event happens**.
- **What is a Callback**
  - A **callback** is a mechanism where **you give your function (URL/code) to another system**, and that system will **call you back later when something happens**.
- **Features**
  - It is **event-driven, one-way, push-based communication**.
  - Unlike API polling, you **do NOT repeatedly call API to check updates**.
  - Instead, the provider **pushes data automatically** to your endpoint.



55

## Comparison

www.cognixia.com

56

## Comparison



Feature / Category	Message Queue	Pub/Sub	Event Streaming	Webhook
Purpose	Reliable <b>task processing</b> with guaranteed delivery	Real-time <b>event broadcasting</b> to multiple subscribers	<b>Continuous real-time stream processing &amp; analytics</b>	<b>Notify external system</b> when an event happens
Communication	Async, Push or Pull	Async, Push	Async, Push or Pull	Async, Push
Message Delivery	One-to-One	One-to-Many	One-to-Many (Replayable)	One-to-One or One-to-Many
Real-Time	Near real-time	Real-time	Real-time (continuous)	Real-time
Use Case	Background job processing	Notifications & broadcasting	Analytics, tracking, IoT, logging	Payment updates, CRM integration
Industry Example	E-commerce Order Processing	WhatsApp broadcast, stock alerts	Fraud detection, GPS tracking	Stripe payment callback

[www.cognixia.com](http://www.cognixia.com)

57

## Database Design




[www.cognixia.com](http://www.cognixia.com)

58

29

# What is Database Design

- **Database Design** is the **process** of structuring the data and **defining the relationships** between different **data elements** to support the **business needs** of an application.
- The goal is to **produce a model** that ensures data **integrity**, minimizes **redundancy**, and optimizes for **performance (speed of queries)**.
- The fundamental decision in database design is choosing the underlying data model, primarily falling into **two major categories**:
  - Relational
  - Non-Relational (NoSQL)



[www.cognixia.com](http://www.cognixia.com)

59

## Relational Databases (RDBMS)

[www.cognixia.com](http://www.cognixia.com)

60

# Relational Databases (RDBMS)



- Relational databases **store data in tables** with **fixed schemas** (columns) and **rows**.
- They are based on the **relational model** and rely heavily on structured **relationships defined by keys** (*primary and foreign*).
- **Key Characteristics**
  - **Schema**
    - **Fixed and Structured** (*predefined columns and data types*).
    - Changes require migrations.
  - **Relationships**
    - Defined through **Foreign Keys** and **Joins** (e.g., *a customer is linked to their orders*).
  - **Transactions**
    - Adheres to **ACID properties** (*Atomicity, Consistency, Isolation, Durability*), making them ideal for **financial transactions** and systems requiring **high data integrity**.
  - **Scaling**
    - **Primarily scales vertically** (*adding more CPU/RAM to a single server*), though techniques like sharding allow for **horizontal scaling**.

[www.cognixia.com](http://www.cognixia.com)

61

## RDBMS ... continue



- **When to Choose RDBMS**

Condition	Choose SQL?
High data consistency required	<b>Yes</b>
Complex relationships	<b>Yes</b>
Transaction rollback needed	<b>Yes</b>
Scalability needed	<b>Hard</b> via vertical scaling
Unstructured data	<b>No</b>

- **Popular Databases:** MySQL, PostgreSQL, Oracle, SQL Server, MariaDB

[www.cognixia.com](http://www.cognixia.com)

62



## Non-Relational (NoSQL)

[www.cognixia.com](http://www.cognixia.com)

63

## Non-Relational Approaches (NoSQL)

- NoSQL databases encompass various **data models that deviate from the traditional table structure**.
- NoSQL databases store **unstructured, semi-structured, or high-volume data** using flexible data models.
- Focuses on **speed, scalability** and **schema flexibility**.
  - **No fixed schema**
  - **Horizontally scalable (distributed)**

[www.cognixia.com](http://www.cognixia.com)

64

# NoSQL | Types



Type	Data Structure	Best Use Case	Industry Example
Document Store	Stores <b>data</b> as flexible, semi-structured <b>JSON documents</b> .	<b>Content Management, Catalogs, Profiles:</b> Data structure changes frequently, and data is read/written as a single unit.	<b>MongoDB, Couchbase</b> used for storing e-commerce product catalogs.
Key-Value Store	Simple <b>map</b> of a <b>unique key to a value</b> (which can be any object).	<b>Caching, Session Management, Real-Time Data:</b> Requires extremely fast, non-complex lookups.	<b>Redis, Memcached</b> used for storing user session tokens.
Wide-Column Store	<b>Stores data in tables</b> , but columns can vary from row to row; optimized for high availability and big data.	<b>Time-Series Data, Big Data Analytics, High-Volume Logging.</b>	<b>Cassandra, HBase</b> used for logging data from millions of IoT devices.
Graph Database	<b>Stores data as nodes</b> (entities) and <b>edges</b> (relationships).	<b>Social Networks, Recommendation Engines, Fraud Detection:</b> Relationships are the most important part of the data.	<b>Neo4j</b> used for storing a social graph (who follows whom).

[www.cognixia.com](http://www.cognixia.com)

65

# NoSQL | Horizontal Scaling



- In NoSQL databases, **horizontal scaling** is achieved primarily through **sharding** and **replication**.
  - **Sharding (Data Partitioning)**
    - Sharding is the process of splitting the entire dataset into smaller, independent chunks called **shards** or **partitions**.
    - **Mechanism**
      - Each shard is placed on a **separate server (node)**.
      - When a **new server** is added to the cluster, the **system automatically redistributes** some **data** from the **existing nodes** to the **new node**.
    - **Result**
      - A **single query** only needs to hit the **server holding the relevant partition**, drastically improving performance under high load.
  - **Replication (High Availability)**
    - Replication ensures that **multiple copies** of the **data** exist across different nodes.
    - **Mechanism**
      - Even though the **data** is sharded, each piece of **data** is typically copied to several other **nodes**.
    - **Result**
      - If a **node fails**, the system can immediately **retrieve the data** from a **replica** on another active node, ensuring high availability and fault tolerance.
  - **Distributed Operations**
    - When an **application** needs to **read or write** data, the **request goes** to the **cluster**, and the system's **distribution coordinator** knows exactly **which nodes hold** the required **data** (or its **replica**) and **directs the request** accordingly.

[www.cognixia.com](http://www.cognixia.com)

66

# NoSQL | Consistency



- **Eventual Consistency**
  - In a distributed NoSQL system, a **write operation** must be **replicated across multiple nodes** (servers) to ensure high availability and durability.
  - The **consistency issue arises** because the **system allows clients to read data before all replicas have been successfully updated.**
  - **The Issue**
    - **Stale Reads**
      - After a **successful write to one node**, a **client reading the data from a different, yet-to-be-updated node** will receive the **old (*stale*) data**.
    - **Conflict**
      - If multiple clients write to the **same data on different nodes simultaneously**, the **system has a conflict and must determine which write wins**.

[www.cognixia.com](http://www.cognixia.com)

67

## Distributed Relational Database



[www.cognixia.com](http://www.cognixia.com)

68

# What is Distributed Relational Database



- **Distributed Relational Database (DRDB)**, often referred to as **Distributed SQL**, is a database system that **combines** the **relational data model** (*tables, schemas, SQL*) and **ACID transactional guarantees** of a traditional RDBMS with the horizontal scalability, fault tolerance, and global distribution of distributed systems.
- Unlike a **traditional database** that runs on a **single server**, a **DRDB spreads** (*or shards*) and **replicates** its **data across multiple independent servers** (*nodes*) connected by a network.
- This makes the system appear as a single, logical database to the application, while the distributed architecture handles the complexities of data placement, synchronization, and query execution across the cluster.

[www.cognixia.com](http://www.cognixia.com)

69

# Working of Distributed Relational Database



- Key mechanisms that enable a DRDB to function:
  - **Data Sharding (Partitioning)**
    - The database automatically breaks up large **tables** into smaller, manageable **chunks (shards)** and distributes them across different **nodes** in the cluster.
    - This allows for **parallel processing of queries**, improving performance for massive datasets.
  - **Replication**
    - Copies of data (*replicas*) are stored on multiple nodes.
    - If one node fails, the system automatically redirects traffic to a replica, ensuring high availability and fault tolerance.
  - **Distributed ACID Transactions**
    - DRDBs use sophisticated consensus algorithms (like Raft or Paxos) to ensure that transactions maintain **ACID** (*Atomicity, Consistency, Isolation, Durability*) properties, even when a single transaction involves data spread across several nodes.
  - **Horizontal Scalability**
    - Performance and capacity can be increased simply by adding more nodes to the cluster, rather than upgrading the hardware of a single server (*vertical scaling*).

[www.cognixia.com](http://www.cognixia.com)

70



# DNS

[www.cognixia.com](http://www.cognixia.com)

71

## What is DNS



- The **Domain Name System (DNS)** is a core internet service that translates human-readable domain names into computer-readable IP addresses.
- Essentially, it acts as the "**phonebook of the internet**," allowing users to connect to websites and services without memorizing complex numerical addresses.
- DNS is not just about **translating web addresses**; it is a critical component for modern, scalable, and resilient application design.
- **Connection Translation (The Core Function)**
  - **Problem**
    - ✓ Servers are addressed using IP addresses (e.g., *IPv4 or IPv6*) which can change, especially in cloud environments.
    - ✓ **Users only remember domain names.**
  - **DNS Solution**
    - ✓ It provides the mechanism for a client (*web browser, mobile app, etc.*) to query the IP address of the target server.
    - ✓ This decouples the service's location from its name.

[www.cognixia.com](http://www.cognixia.com)

72

# DNS Routing



- DNS routing, often simply called **DNS-based traffic steering** or **load balancing**, is the practice of **using the Domain Name System (DNS)** to intelligently **direct a user's connection** request to a **specific server, IP address, or endpoint** based on criteria **beyond simple name-to-address translation**.
- It's a crucial **technique** in modern system architecture for achieving **high availability, fault tolerance, low latency, and geographic compliance**.

Strategy	Routing Logic	Primary Result & Goal	Use Case Example
Geo-based Routing	Analyzes the <b>source IP address</b> to determine the user's geographical location.	Returns the IP address of the server or CDN endpoint that is <b>geographically closest</b> to the user. <b>Goal: Low Latency &amp; Compliance</b> .	Directing users from Asia to the Singapore data center for faster load times.
Latency-based Routing	Constantly measures the <b>actual network latency</b> between the user's region and various server endpoints.	Returns the IP address of the endpoint that currently offers the <b>lowest measured latency</b> , regardless of physical distance. <b>Goal: Optimal Performance</b> .	Sending a user to a slightly further but less congested server that offers a quicker response time.
Failover Routing	Continuously runs <b>health checks</b> against the application servers to monitor their operational status.	If the primary server fails a health check, its IP is removed, and the IP of the designated <b>secondary (failover)</b> server is returned. <b>Goal: High Availability &amp; Resilience</b> .	Automatically switching all traffic from Data Center A to Data Center B when A suffers a power outage.
Weighted Routing	Assigns a numerical <b>weight</b> (e.g., 80/20, 50/50) to each available IP address based on desired traffic distribution.	Distributes traffic proportionally according to the defined weights. <b>Goal: Controlled Deployment &amp; Load Balancing</b> .	Sending 10% of users to a new version of the application (Canary Deployment) before rolling it out widely.

73

# Storage


[www.cognixia.com](http://www.cognixia.com)


74

# What is Storage

- Storage in system architecture refers to the components and technologies responsible for digitally **recording and retaining data** for immediate or future use.
- It is a fundamental element of any computing system, determining how much data can be kept, how quickly it can be accessed, and how reliably it is protected.
- The architect's view focuses on three main dimensions: **Performance/Access, Sharing/Networking, and Durability/Cost.**

## DATA STORAGE TYPES: Block, File, Object

Organizing Data for Different Needs



Benefit: Choose the right storage type to optimize for performance, cost, and scale ✨

[www.cognixia.com](http://www.cognixia.com)

75

# Performance and Access Types

- Select storage based on the required speed and the fundamental way data is accessed.

Access Type	Description	Key Metric	Typical Technology	Use Case Example
Block Storage	<ul style="list-style-type: none"> <li>Data is stored in <b>fixed-size blocks</b>, each with a unique address, allowing for highly efficient, direct read/write access.</li> <li>The operating system sees it as a raw disk.</li> </ul>	IOPS (I/O Operations Per Second) and Latency	SAN, Local SSD/NVMe	Databases, Virtual Machines (VMs), boot volumes.
File Storage	<ul style="list-style-type: none"> <li>Data is stored as <b>files and directories</b>, complete with metadata (permissions, last modified date).</li> <li>Access is hierarchical.</li> </ul>	Throughput and Latency	NAS (NFS/SMB)	User home directories, centralized file shares, content repositories.
Object Storage	<ul style="list-style-type: none"> <li>Data is stored as self-contained "objects" (files plus all their metadata) in a flat, massive pool, accessed via HTTP APIs.</li> <li>Highly scalable but slower.</li> </ul>	Scalability and Durability	Cloud Storage (S3, GCS)	Data lakes, backups, static web content, media libraries.

76



[www.cognixia.com](http://www.cognixia.com)

77