



## Scaling Systems for Growth

By: Ahmad



1



## Scalability Pattern

[www.cognixia.com](http://www.cognixia.com)

2

# Scalability Patterns



- Scalability Patterns are reusable design methods that help a system handle more users, more traffic, and more data without slowing down or failing.
- They are proven ways to make software grow smoothly as demand increases.
- Why do we need them?
  - Because when systems get bigger, they face:
    - Performance issues
    - High traffic spikes
    - Database overload
    - Slow response times
- Scalability patterns help architects solve these problems efficiently and predictably.



www.cognixia.com

3

# Core Scalability Patterns



- Scalability patterns focus on distributing load and decoupling system components.
- Here are the most fundamental ones:
  - **Load Balancing (Traffic Distribution)**
    - Using a specialized device or software (*Load Balancer*) to distribute incoming network traffic across multiple application servers.
    - **Benefit**
      - ✓ Prevents any single server from becoming overwhelmed, ensuring maximum throughput and high availability.
      - ✓ If one server fails, the load balancer routes traffic to the healthy ones.
  - **Caching (Speed and Database Offload)**
    - Storing the results of frequently requested or expensive operations (like a complex database query or a rendered webpage) in a faster, temporary storage layer (like Redis or Memcached).
    - **Benefit**
      - ✓ Reduces latency for users and significantly reduces the load on the slower primary data store (database).
  - **Sharding or Partitioning (Data Distribution)**
    - The process of splitting a single, large logical database or dataset into multiple, smaller, independent databases, called shards or partitions, each running on its own server.
    - **Benefit**
      - ✓ Allows the database to scale horizontally, overcoming the storage and processing limits of a single machine.

www.cognixia.com

4

## Core Scalability Patterns ... continue



- **Replication (Redundancy and Read Scaling)**
  - Creating and maintaining **multiple copies** of **data** or services **across different servers**.
  - **Benefit**
    - ✓ **Improves fault tolerance** (*if one copy fails, others exist*) and read scalability by **allowing read requests** to be **served** by the **copies** (*known as read replicas*) **instead of overloading** the primary **server**.
  
- **Auto-Scaling (Elasticity)**
  - Automatically **adjusting** the **number of computing resources** (servers) in use based on the current load or demand.
  - **Benefit**
    - ✓ **Ensures optimal performance during peak hours** by scaling up and saves cost during quiet hours by scaling down.

[www.cognixia.com](http://www.cognixia.com)

5

## Scaling



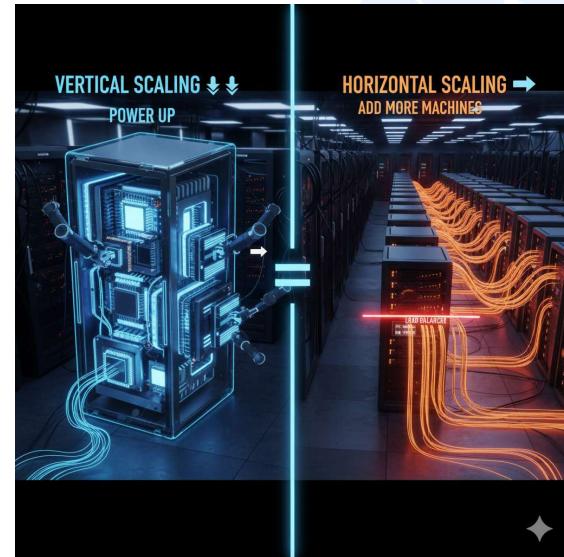
[www.cognixia.com](http://www.cognixia.com)

6

# What is Scaling



- Scaling is the ability of an IT system or business process to **handle an increasing amount of work** (more users, data, or transactions) by adding resources efficiently.
- It ensures that as demand grows, the system remains fast, available, and reliable.
- Auto scaling involves **two main concepts**:
  - Scaling Types** (the direction of change)
  - Scaling Methods** (the trigger for the change).
- There are **two main ways** to increase a system's capacity in a technical context:
  - Vertical (Scaling Up)**
  - Horizontal (Scaling Out)**



[www.cognixia.com](http://www.cognixia.com)

7

## Vertical Scaling (Scale-Up)



- Vertical scaling, or **scaling up**, means **increasing the capacity** of a **single machine**.
- You **add more resources**—like CPU, RAM, or faster storage (SSD)—to an **existing server**.
- Limit:** It's **limited** by the **maximum hardware specifications** available for a **single physical machine**.
- Availability:** Creates a **Single Point of Failure** (SPOF).
- Simplicity:** It is generally **simpler to manage** since you only have one machine to maintain and don't need complex distributed logic.
- Example**
  - Upgrading your laptop from **8 GB** → **32 GB** RAM to run more apps.
- When to Use Vertical Scaling**
  - Application** is **monolithic** or tightly coupled
  - You **cannot add multiple servers**

[www.cognixia.com](http://www.cognixia.com)

8

## Horizontal Scaling (Scaling Out)



- Horizontal scaling, or **scaling out**, means **increasing capacity** by **adding more machines** to the system.
- You **add more servers** (or virtual instances/containers) and **distribute the workload among them**.
- **Limit:** Virtually **limitless** (you can theoretically keep adding servers).
- **Availability/Fault Tolerance:** If **one server fails**, the load balancer **automatically directs traffic** to the **remaining healthy** servers, meaning no downtime for users.
- **Complexity:** Requires architectural changes to ensure the **application is stateless** (so any server can handle any user request) and requires managing **data consistency** across **multiple nodes**.
- **When to Use Horizontal Scaling**
  - Microservices-based architecture
  - **Stateless applications**
  - **High traffic web and mobile apps**
  - **Dynamic workloads**

[www.cognixia.com](http://www.cognixia.com)

9

## Auto-Scaling



- Its **main purpose** is to ensure that **applications maintain consistent performance and availability** by always having the necessary resources, while **simultaneously optimizing costs** by **avoiding the over-provisioning** of resources **during low-demand periods**.
- The goal is to **always match capacity with demand**:
  - **When traffic spikes:** Automatically add resources to prevent slowdowns.
  - **When traffic drops:** Automatically remove resources to save money.
- In a **public cloud** (like AWS or Azure), **auto scaling** is a **service they provide and manage**.
- **On-premises** auto scaling is when you **build this same capability** in your **own data center**.
- It's **more complex** because you have to manage all the moving parts.
- The **most common way** to do this today is by using **Kubernetes** (often called *K8s*), which is the software platform used to manage application containers.

[www.cognixia.com](http://www.cognixia.com)

10

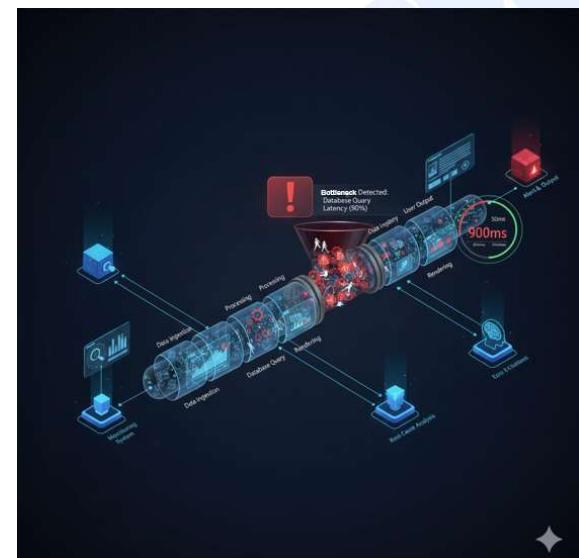
# Performance Analysis

[www.cognixia.com](http://www.cognixia.com)

11

## What is Performance Analysis

- Performance Analysis is the **process of measuring how well a system performs** and finding **what slows it down**.
- Its **primary goal** is to **ensure the system meets required performance objectives** (like fast response times) and to **identify and resolve** the **bottlenecks** that limit its capacity.
- It helps identify:
  - Bottlenecks
  - Latency issues
  - Slow queries
  - High CPU/RAM usage
  - Network delays
  - Scaling problems
- **Goal:** Make the **system faster, stable, and efficient under real load**.

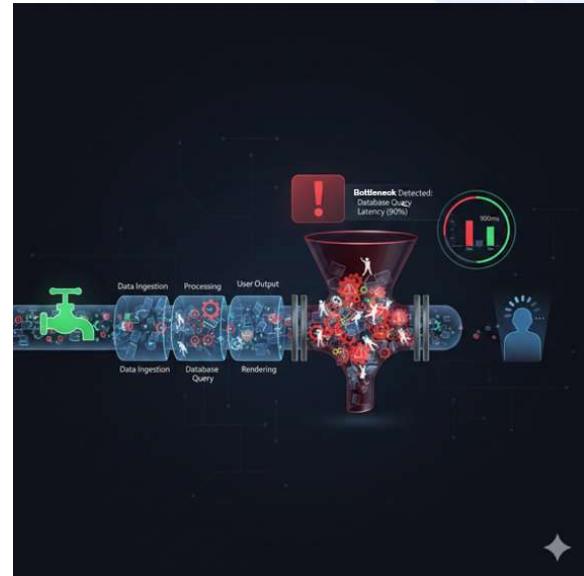

[www.cognixia.com](http://www.cognixia.com)

12

## Bottlenecks Analysis



- The core objective of this **analysis** is to find and **eliminate bottlenecks**—the **components** that **constrain** the **overall throughput** or **latency** of the system.
- A **bottleneck** is any component that slows down the entire system — like the slowest person in an assembly line.



[www.cognixia.com](http://www.cognixia.com)

13

## Detecting Bottlenecks



Layer / Component	Symptoms (What You See)	How to Detect (Tools & Techniques)
Application Layer	<b>Slow APIs, high CPU, hangs</b>	<b>APM tools</b> (Datadog, NewRelic), <b>CPU/Memory profiling</b> , <b>Logs, Traces</b>
Database Layer	<b>Slow queries, locks, DB CPU 80–100%</b>	Query execution plans, <b>DB monitoring</b> dashboards, <b>slow query logs</b>
Storage/I/O Layer	<b>File uploads slow, high I/O wait</b>	<b>Disk IOPS metrics</b> , CloudWatch, Storage performance insights
Network Layer	<b>High latency, packet drops</b>	<b>Ping tests, traceroute</b> , <b>VPC Flow Logs</b> , Network dashboards
Messaging Queue	<b>Queue backlog increasing</b>	<b>Queue depth metrics</b> , consumer lag monitoring, CloudWatch/Prometheus
Third-Party APIs	<b>External call latency spikes</b>	<b>API Gateway logs</b> , circuit-breaker metrics, distributed tracing
Front-End/UI	<b>Slow page load, high Time-to-First-Byte</b>	<b>Web Vitals</b> , Chrome DevTools, Lighthouse, RUM tools

[www.cognixia.com](http://www.cognixia.com)

14

## Fixing Bottlenecks



Layer / Component	Fix / Optimization Actions
Application Layer	Add caching (Redis), optimize code loops, async processing, add Auto Scaling, refactor heavy logic
Database Layer	Add indexes, optimize queries, add read replicas, increase connection pools, sharding/partitioning
Storage/I/O Layer	Move to SSD, increase provisioned IOPS, use object storage (S3), enable caching/CDN
Network Layer	Reduce cross-region calls, use private links, increase bandwidth, enable API caching
Messaging Queue	Add more consumers, increase batch size, change visibility timeout, auto-scale workers
Third-Party APIs	Add circuit breaker, retry with backoff, local caching, queue requests asynchronously
Front-End/UI	Compress images, minify JS/CSS, enable CDN, lazy load assets, optimize bundle size
Microservices	Reduce network hops, optimize service boundaries, tune gRPC/HTTP calls, add bulkheads

15



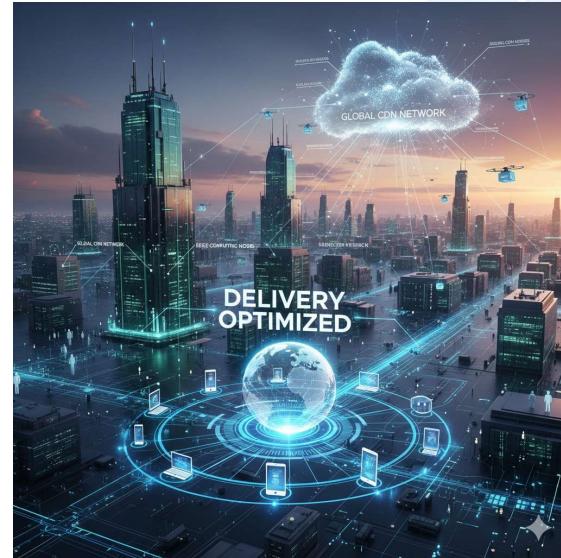
## Optimizing Delivery

[www.cognixia.com](http://www.cognixia.com)

16

# Optimizing Delivery

- Optimizing delivery with **Content Delivery Networks (CDNs)** and **Edge Computing** involves **moving data and computation geographically closer to the end-users** to dramatically reduce latency, decrease network congestion, and improve the overall user experience.
- **Content Delivery Network (CDN)**
  - A CDN is a **geographically distributed network of proxy servers** and data centers.
  - It focuses primarily on the **delivery of content**.

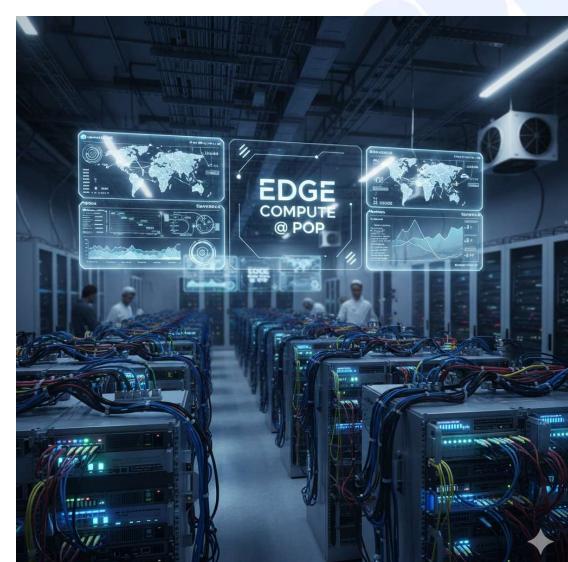


[www.cognixia.com](http://www.cognixia.com)

17

# Edge Computing

- Edge Computing is an **architectural approach** that **extends** the concept of a **CDN** by **moving application logic and computation—not just cached content—to the edge of the network**, even closer to the user or the data source.
- **Computation at the Edge**
  - Unlike a traditional CDN which primarily caches static files, **Edge Computing runs small applications** or functions (*known as "serverless functions" or "Edge functions"*) **on the PoPs**.
- **Preprocessing Data**
  - It can **handle tasks that require low latency**, such as data validation, basic authentication, formatting data, or simple business logic, before the request ever reaches the main cloud data center.



[www.cognixia.com](http://www.cognixia.com)

18

## Edge Computing ... continue



- **Key Benefits**
  - **Ultra-Low Latency**
    - Processing is done immediately **at the edge**, making it vital for **time-sensitive applications**.
  - **Bandwidth Efficiency**
    - Edge devices can filter, **aggregate**, and **analyze data locally**, sending only relevant information back to the central cloud, saving bandwidth.
  - **Enhanced Security**
    - Edge functions can be used to **scrub incoming requests** for **malicious traffic** before they reach the core infrastructure.
- **Industry Example**
  - **Industrial IoT (Internet of Things) Manufacturing**
    - Sensors on a **factory floor** generate **massive amounts of data in real-time**.
    - Instead of sending all that **raw data** to a **central cloud** (which would be slow and expensive), an **edge gateway** in the **factory processes** the **data instantly** to detect equipment anomalies.
    - This immediate, **local processing** allows for **rapid shut-down** of a **malfunctioning machine** (critical action) long **before** the **central cloud could ever receive and analyze the alert**.

[www.cognixia.com](http://www.cognixia.com)

19



[www.cognixia.com](http://www.cognixia.com)

20