



Observability and Maintenance

By - Ahmed



www.cognixia.com

1

What is Observability

- Observability is the ability to **understand the internal state** of a complex software system solely by examining its **external outputs** (*telemetry data*).
- In essence, an **observable system** gives you the **flexibility** to explore and ask arbitrary questions **about its behavior**—questions you didn't even know to ask when the system was deployed.
- It **moves beyond knowing that something is wrong** (traditional monitoring) to **understanding why** it is wrong and **how** to fix it.
- **Tools:** ELK Stack (Elasticsearch, Logstash, Kibana), Splunk, Jaeger, Zipkin, etc.



www.cognixia.com

2

Observability vs Monitoring



- The terms are often **used interchangeably**, but there is a **crucial distinction** that reflects the **complexity of modern, distributed architectures** (like microservices):
 - Monitoring (The "Known Unknowns")**
 - Tells you **when something is wrong** and **what the basic problem is** (e.g., "CPU utilization is at 95%", or "The error rate is above 1%").
 - Monitoring relies on **predefined metrics** and dashboards.
 - Observability (The "Unknown Unknowns")**
 - Tells you **why the problem is happening** and **how to fix it**.
 - It allows engineers to **drill down** into the **data** to **trace** a **fault** through **hundreds of services** and **discover** the **root cause**, even for **failures** they **hadn't anticipated** or **pre-configured** an **alert** for.

www.cognixia.com

3

Pillars of Observability



- Observability relies on the continuous collection and correlation of three primary types of telemetry data:
- Metrics (The "What?")**
 - Metrics** are **quantitative, numerical measurements** of a service or system captured over time.
 - Format**
 - Time-series data** (`timestamp, value, name/tag`).
 - Purpose**
 - Provide high-level insights** into **resource utilization** and **service health**.
 - Excellent for setting alerts and trending over time.
 - Examples**
 - CPU usage, memory consumption, request latency, requests per second, error rate.**



4

Pillars of Observability ... continue



- **Logs (The "Why?")**
 - Logs are discrete, time-stamped text records of events that occurred within an **application or system**.
 - **Format**
 - *Text (structured JSON or unstructured plain text).*
 - **Purpose**
 - **Provide granular context for debugging.**
 - Logs are used to **pinpoint the exact sequence of events** that led to an error.
 - **Examples**
 - User login successful, database query failed, function entered, exception thrown.
- **Traces (The "Where?")**
 - Traces track the **end-to-end journey** of a **single request** as it **flows** through **all services, components, and databases** in a distributed system.
 - **Format**
 - A collection of **timed spans**, each representing a **segment of work done** in a service.
 - **Purpose**
 - **Provide visibility** into the **relationships and dependencies** between services.
 - Essential for identifying latency **bottlenecks** in microservice architectures.
 - **Example**
 - A request hits the **API Gateway** → **authenticates** with the **Auth Service** → **queries** the **User Profile Service** → **queries** the **Product Catalog Database** → **returns a response**.
 - The trace shows how long each step took.

www.cognixia.com

5

Proactive Health Check



www.cognixia.com

6



Proactive Health Checks and Alerting

- **Proactive System Health Checks and Alerting** is a core strategy in modern IT operations that moves system management from being **reactive** (*fixing things after they break*) to **preventative** (*identifying and resolving potential problems before they impact users*).
- It is achieved by **continuously measuring key system indicators** against established **performance baselines** and immediately notifying responsible teams when a **threshold** is breached.



7



Tools for Health Checks

- A comprehensive strategy requires tools that can **collect metrics**, **visualize** them, and manage the **alerting lifecycle**.

Category	Open-Source Tools	Commercial/Managed Platforms
Metrics & Visualization	Prometheus (Metrics Collection & Storage), Grafana (Dashboards)	Datadog , Dynatrace, New Relic
Infrastructure Monitoring	Nagios , Zabbix , Icinga	SolarWinds, ManageEngine OpManager
Incident Management	Alertmanager (for Prometheus)	PagerDuty , Opsgenie (Atlassian), Squadcast
Log Management	ELK Stack (Elasticsearch, Logstash, Kibana)	Splunk , Datadog Logs

www.cognixia.com

8



Service Level

www.cognixia.com

9

Service Level

- A Service Level refers to the specific, measurable **performance standard** or target that is agreed upon for the **delivery** of a **service**.
- It essentially answers the question: "**How well and how quickly will this service perform?**"
- **The Interconnected Relationship**
 - Think of them as a **hierarchy of commitment**:
 - **SLI: THE MEASUREMENT** (*The actual percentage of uptime, e.g., 99.95%*)
 - **SLO: THE TARGET** (*We aim for 99.9% Uptime*)
 - **SLA: THE COMMITMENT** (*If we fail to meet the 99.9% target, we pay you a penalty*)



www.cognixia.com

10

Service Level | SLI



- **Service Level Indicators (SLIs)**

- SLI is a **numerical measure** that tells **how reliable** the **system is**. The **raw data points** that tell you *what* the service is doing.
- A **quantitative measure** of a service's performance (e.g., *99.99% of requests returned in under 100ms*). This detects risk.
- SLIs **measure customer-impacting metrics**, such as:

SLI Type	What It Measures	Example Metrics
Availability	Service uptime, success response rate	% of successful API calls (2xx)
Latency	Speed of response	% of requests completed under 200 ms
Error Rate	Failed or erroneous transactions	5xx errors, failed login attempts
Throughput / Traffic	Requests handled per second	API requests/second, trades/sec
Saturation / Capacity	Resource exhaustion levels	CPU, memory, queue usage
User Experience (UX)	Real user performance	App crash rate, page load time
Business SLI	Business/user success	% successful trades, payment success rate

11

Service Level | SLOs



- **Service Level Objectives (SLOs)**

- The agreed-upon **target reliability percentage** (e.g., *99.95% uptime*), based on SLI data.
- The explicit, **internal goal** for the reliability of a service. The SLO is an **internal contract** between the SRE team and the Development (Product) team.
- **Example**
 - The payment API should be 99.95% available over 30 days.
 - 95% of trades should complete within 300ms.
- It tells **how reliable** the service **SHOULD be**, not to be perfect, but **good enough for customers** at a **reasonable cost**.

SLO Type	What It Measures	Example Metric	Example SLO Statement	Used For
Availability SLO	% of time service is available	Uptime %, Error rate	Service will be available 99.9% of the time per month	Service reliability, uptime commitments
Latency / Performance SLO	Speed of response	Response time (ms), Page load time	95% of API calls should respond under 200ms	User experience, app responsiveness
Throughput SLO	Volume of successful requests per time	Requests/sec, Transactions per minute	System should handle 5000 orders per minute	Scalability, system capacity
Error Rate SLO	Frequency of failed requests	% failed requests	Less than 0.5% of API requests can fail in a week	Quality of service, reliability
Durability SLO	Protection of data from loss	Data loss %	Data durability of 99.999999999% (11 nines) per year	Data storage, backup, cloud services
Reliability SLO	Successful task completion	Successful job completion rate	99.9% of scheduled jobs should finish successfully	Background jobs, workflows

12

Service Level | SLA



- **Service Level Agreements (SLA)**
 - **Contract**
 - The **external, legal agreement** with the **customer**, outlining **penalties** for **failure to meet** the **SLO** (e.g., **99.5%** with a **10% credit** if missed).
 - **Buffer**
 - SLOs are typically set **tighter** than the **external SLA** to provide a **safety margin**, ensuring that if the team misses the **internal target** (the **SLO**), they still have a chance to fix it before breaching the **external legal contract** (the **SLA**).

SLA Type	What It Ensures	Real-Life Example
Availability SLA	Uptime guarantee	Server SLA: 99.99%
Support SLA	Response/resolution time	P1 issue responded in 15 min
Performance SLA	Response speed guarantee	95% of transactions < 200ms
Data Protection SLA	Backup, retention, durability	99.999999999% (11 nines) data durability
Compliance SLA	Security, regulation, audit	GDPR, PCI-DSS, RBI guidelines
Penalty-based SLA	Refund/credit if breached	Refunds 25% monthly bill

www.cognixia.com

13

Pipelines




www.cognixia.com

14

What is Pipelines



- CI/CD Pipelines are automated processes that take code changes from a developer's machine and reliably deliver them to production.
- They are the backbone of modern software development, enabling rapid, frequent, and reliable software releases.
- The acronym stands for Continuous Integration (CI) and Continuous Delivery/Deployment (CD).



www.cognixia.com

15

Phases of the Pipeline



- **Continuous Integration (CI)**
 - CI focuses on automatically merging developer code changes into a central repository frequently.
 - **Action**
 - A developer commits code.
 - **Process**
 - The pipeline automatically performs:
 - ✓ **Building**
 - ❖ Compiling the code into an executable artifact (e.g., a JAR file, Docker image).
 - ✓ **Testing**
 - ❖ Running unit tests, integration tests, and static code analysis to immediately catch bugs or security flaws.
 - **Outcome**
 - If tests pass, a stable, ready-to-deploy artifact is created and stored (e.g., in a container registry).
 - **Continuous Delivery / Deployment (CD)**
 - CD focuses on automating the release of the validated artifact to various environments (staging, production).
 - **Continuous Delivery**
 - The tested artifact is delivered to a central repository (e.g., Docker registry) and is ready for manual, one-click deployment to production.
 - A human decision is required for the final production push.
 - **Continuous Deployment**
 - The tested artifact is automatically deployed to the production environment without any human intervention.
 - Automation is complete.

www.cognixia.com

16

Safe Rollbacks



- The **pipeline** is **designed** not just for forward deployment but also for **safely reversing a change** if a critical error is detected in production.
- **Automated Verification**
 - After deployment, the pipeline runs post-deployment health checks (e.g., synthetic transactions, latency checks) to confirm the new version is operating correctly.
- **Rollback Trigger**
 - If the **health checks** fail or if monitoring/observability tools detect a critical performance degradation (e.g., a spike in error rates), the pipeline's failure path is automatically executed.
- **Rollback Mechanism**
 - The most common safe rollback strategy is to re-deploy the immediately preceding, known-good version of the application artifact.
 - Because the pipeline archives all successful artifacts, reverting is simply another automated deployment process.
 - ✓ **Canary/Blue-Green Rollbacks**
 - ❖ Modern CD strategies use techniques like **Blue/Green** (maintaining the old version alongside the new) or **Canary Releases** (rolling out the new version to a small subset of users).
 - ❖ If the new version fails, traffic is instantly shifted back to the stable, **old environment** with minimal impact on users.
 - ❖ This ensures the system maintains high availability even when a new version is failing. www.cognixia.com

17

Q & A

concepts still unclear?

Thank you for attending

www.cognixia.com

18



19