

ForgeSync "Global Industry 4.0"

Project Sector: Precision High-Tech Manufacturing (Automotive & Aerospace)

1. Project Overview: The Digital Nervous System

The Vision: "ForgeSync Core" is a next-generation **On-Premises-First** architecture. Unlike standard IT systems, **manufacturing cannot rely on a constant internet connection**. If the **connection** to the **cloud drops**, the **factory cannot stop**.

The Core Problem: The "Siloed" Factory

ForgeSync's 4 Global Mega-Factories (*Detroit, Stuttgart, Shanghai, São Paulo*) are currently "islands."

- **The 1% OEE Problem:** **Overall Equipment Effectiveness** is stuck at **68%**. Because data isn't synchronized, **managers don't know why machines** are **slowing** down **until the shift ends** (*This means 32% of our time and money is being wasted because machines are breaking down, running too slowly, or making broken parts*).
- **The Value:** **Moving** from 68% **to 69% OEE** is worth **\$45 Million/year** in **saved** materials, reduced energy, and higher output.
- **The Constraint:** All **production-critical logic** (*Machine Control, Safety, Quality Vision*) **must** run **On-Premises** at the Factory Edge to **ensure sub-10ms latency** and **99.99 uptime** during **network outages**.

2. Business Scenario & Operational Environment

The "On-Premises" Hierarchy:

Your **architecture must follow** the **industrial standard** where layers are **separated by function** and **physical location**:

Layer	Location	Function	Requirement
Layer 0-2 (Shop Floor)	On-Premises	PLCs, Sensors, Robotic Arms.	Ultra-Low Latency: <10ms response time.
Layer 3 (Factory Edge)	On-Premises	Manufacturing Execution System (MES), Local Data Historian.	High Autonomy: Must run for 48 hours without internet.
Layer 4 (Corporate)	Global HQ	ERP Integration, Global OEE Analytics, Long-term AI Training.	Global Visibility: Aggregated data for executive decision-making.

The "Forge" (Hardware Details):

1. **The Robotic Arm (The Worker):** A large **mechanical arm**. **If it moves too fast, it vibrates.** If it rubs against metal, it screeches. **If it works too hard, it gets hot.**
2. **The PLC (The Machine's Local Controller):** A rugged, **mini-computer inside** the machine. It executes "Move" or "Stop" commands. Your **software must tell** the PLC when to "STOP" instantly.
3. **The Sensors (The Senses)**

A **sensor** is a **tiny device** that **converts** a **physical** "feeling" into a **digital number**. You will **monitor 4 main sensors**:

Sensor Name	What it feels	Digital Data	The "Danger" Sign
Vibration	The "shake" of the motor	Hertz (Hz)	<ul style="list-style-type: none"> ○ 12Hz is a happy hum. ○ 85Hz means a gear is loose.
Acoustic	High-pitched sound	Decibels (dB)	<ul style="list-style-type: none"> ○ 70dB is normal. ○ 95dB is a "Grinding" metal sound.
Thermal	Heat	Celsius (°C)	<ul style="list-style-type: none"> ○ 50°C is fine. ○ >80°C means the motor is melting.
Cycle Time	Speed of work	Seconds	<ul style="list-style-type: none"> ○ 2.1s is perfect. ○ 2.5s means the robot is lagging.

Operational Constraints:

1. **Data Sovereignty:** By law, detailed "**Build Logs**" for aerospace parts produced in **Shanghai** or **Stuttgart** must **stay** on the **factory's physical servers for 10 years**. **Only "Summarized Health Stats"** can **leave** the **building**.
2. **The Firehose (Data Load):** 4 Factories x 12 Lines x 250 Sensors.
 - **High-Velocity (Vibration/Acoustic): 100 reports per second** (100 Hz).
 - **Low-Velocity (Thermal/Speed): 1 report per second** (1 Hz).
3. **Safety Interlock:** If the **centralized system sends a "Stop" command**, it must **reach** the machine in **under 5ms**.
4. **Hardware:** **Each factory** has its **own dedicated server room** with **high-availability clusters**.
5. **Latency Constraint (<5ms):** If a **machine** is **"Grinding"** you **cannot send** that **data** to the **Cloud** and **wait** for an **answer**. You **must process** it **locally** in **under 5 milliseconds** to **save** the **robot**.
6. **Autonomy Constraint (100% Uptime):** Even if the **factory loses internet**, the robots and your **software must keep working**.

3. The Assignment: 12 Detailed Tasks

Section 1: Strategic Mapping

- **Assignment 1: Business Vision to Technical Vision.**
- **Assignment 2: Functional & Non-Functional Requirements.**

Section 2: Architectural Selection

- **Assignment 3: Select Paradigm.**
- **Assignment 4: Select Model.**
- **Assignment 5: Select Architecture Style.**
- **Assignment 6: Select Architecture Pattern.**

Section 3: Technical Design & Flow

- **Assignment 7: High-Level Design (HLD).**
- **Assignment 8: Low-Level Design (LLD).**
- **Assignment 9: Component & Service Selection.**
- **Assignment 10: Create 3 ADRs (Architectural Decision Records).**

Section 4: Visualizing the Flow

- **Assignment 11: Create System Flow.**
- **Assignment 12: Final Architecture Picture.**

Model Answer

1. Business Vision to Technical Vision

Business Goal: Increase OEE by 1% to capture \$45 Million/year in value.

Technical Vision: Deploy an On-Premises Digital Twin architecture that provides the "Digital Nervous System" required for high-speed automated decision-making.

Business Driver	Business Goal	Technical Objective (The "How")	Target Metric (The "Number")
Availability	Reduce unplanned machine breakdowns.	Predictive Health. Use Vibration/Thermal sensors to predict failure.	99.99% Availability of critical machines.
Performance	Eliminate "micro-stops" and slow cycles.	Millisecond Tracking. Real-time tracking machine cycle times.	<10ms Data Latency from Sensor to Dashboard.
Quality	Reduce scrap and wasted materials.	Automated Interlocks. Use Acoustics to trigger emergency stops.	<5ms Response Time for automated safety shut-offs.
Resilience	Maintain production during WAN outages.	Edge Autonomy. Keep all "Digital Twin" logic On-Premises.	100% Factory Uptime during Internet/Cloud outages.
Reliability	Ensure data for \$45M savings is accurate.	High-Fidelity Sync. Robust local data storage and synchronization.	Zero Data Loss (RPO=0) for production-critical logs.

2. Functional & Non-Functional Requirements (FRs & NFRs)

- a. **Functional Requirements (FR):** What the "**ForgeSync Core**" system **MUST** do.

ID	Requirement Name	Description
FR-1	Digital Twin Sync	The system shall capture and map physical sensor data (Vibration, Thermal, Acoustic, Cycle Time) to a digital record in real-time.
FR-2	Automated Interlock	The system must automatically send a "Stop" command to the PLC if Thermal levels exceed 80°C or Acoustics detect grinding.
FR-3	Local Dashboarding	The system shall provide a real-time OEE dashboard at the factory floor that updates every minute without requiring internet.

FR-4	Data Buffering	The system shall store all production data locally for up to 48 hours during a WAN/Internet outage.
FR-5	Predictive Alerting	The system shall trigger a "Maintenance Alert" when vibration patterns deviate from the baseline (e.g., rising from 12Hz toward 85Hz).

b. Non-Functional Requirements (NFR): *How the system must perform to protect the \$45 Million value.*

ID	Category	Metric/Target	Reason for North Star
NFR-1	Latency	< 5ms (<i>End-to-End</i>)	Safety: To trigger an emergency stop before machine damage occurs.
NFR-2	Availability	99.99% (<i>High-Availability</i>)	Uptime: Every minute of system downtime causes a drop in OEE.
NFR-3	Resilience	100% On-Prem Autonomy	Reliability: Factory production must not stop if the Cloud/Internet fails.
NFR-4	Data Integrity	Zero Data Loss (<i>RPO=0</i>)	Financials: We cannot calculate the \$45M savings if we lose the production logs.
NFR-5	Scalability	1.2M Events/Sec	Growth: The system must handle the massive "Firehose" of vibration and acoustic data from 4 global factories.

c. Traceability Matrix

- "We have a Functional Requirement to stop the machine if it gets too hot (**FR-2**).
- However, that requirement is useless unless we meet the Non-Functional Requirement of **<5ms Latency (NFR-1)**.
- If the latency is **2 seconds**, the machine is already destroyed by the time the stop command arrives."

d. The Calculation: "The Firehose Math"

To find the scalability requirement, we look at the **Source** (*The Sensors*) and the **Volume** (*The Factories*).

i. The Sensor Count (Density)

- Each of the **4 Global Factories** has **12 Production Lines**.
- Each Production Line has **250+ Sensors**.
- **Total Sensors:** $4 \times 12 \times 250 = 12,000$ sensors globally.
- **Events per second:** 12,000 sensors x 100 signals/sec = **1,200,000**.

ii. Why this NFR is "The Firehose"

In your architecture, this number is a **Stress Test**.

If you design a system that can only handle 10,000 events per second, your "Digital Twin" will be "blind" to 99% of the vibration data.

Without this NFR:

- You miss the "Micro-shake" that leads to a breakdown.
- You lose the **Availability** of the machine.
- You fail the **Metric** (\$45M savings).

3. Select Paradigm

For the ForgeSync platform, the most effective choice is the **Reactive Paradigm** (*specifically using an Event-Driven approach*).

Reactive is a way of writing code that focuses on **Data Streams** and the propagation of change.

Reason	Detail
High Velocity	<ul style="list-style-type: none">○ The system must process 1.2M events/sec (Vibration, Acoustics).○ Traditional "Request-Response" (Imperative) would crash under this load.
Responsiveness	<ul style="list-style-type: none">○ Reactive systems are designed to stay responsive under heavy load by being non-blocking.
Resilience	<ul style="list-style-type: none">○ In a Reactive paradigm, if one component fails, the system stays alive (Isolation).
Real-time Logic	<ul style="list-style-type: none">○ The "Digital Twin" needs to react immediately to a "Vibration Spike" or "Thermal Danger."

How it Works in ForgeSync (The Model)

In this paradigm, we do not "ask" the machine for its temperature every second.

Instead, the machine **"publishes"** its state, and our system **"reacts"** only when something important happens.

- 1) **Produce:** The Physical Sensor (Vibration) produces an "Event."
- 2) **Stream:** The event is pushed into a high-speed stream (e.g., MQTT or Kafka).
- 3) **React:**
 - **Logic A:** If vibration is **12Hz** (Normal) → Simply store it for OEE reports.
 - **Logic B:** If vibration is **85Hz** → (Danger) → Trigger the **Emergency Stop** immediately.

4. Create the Model

a. The Digital Twin Concept

In the "**ForgeSync**" architecture, the Digital Twin is a **live data object** that mirrors the physical state of a machine.

It consists of three parts: **Identity**, **Telemetry** (The Senses), and **Logic** (The Brain).

b. Physical-to-Digital Mapping Table

This table shows how we translate physical reality into digital data for a single **Robotic Milling Arm**.

Physical Entity	Sensor Type	Digital Attribute (Field Name)	Normal Range	Trigger Logic (The Brain)
Machine ID	Asset Tag	twin_id	N/A	Unique Identifier (e.g., FORGE-DET-01)
Motor Health	Vibration (Accelerometer)	vibration_hz	10Hz - 20Hz	If > 85Hz: Harmonic Spike Detected (<i>Alert Maintenance</i>).
Friction Level	Acoustic (Microphone)	noise_db	60dB - 75dB	If > 95dB: Grinding Detected (<i>Trigger Emergency Stop</i>).
Speed/Flow	Cycle Time (Laser)	cycle_time_ms	2000 - 2100ms	If > 2300ms: Performance Lag (<i>Update OEE Dashboard</i>).
Motor Heat	Thermal (Thermocouple)	temp_celsius	40°C - 65°C	If > 80°C: Overheating (<i>Trigger Safety Shutdown</i>).

c. Justification

- **Predictive:** By mapping `vibration_hz`, we catch **Harmonic Spikes** early, preventing a \$500,000 motor failure.
- **Reactive:** By mapping `noise_db`, we detect **Grinding** and stop the machine in **<5ms**, saving the physical machine from destruction.
- **OEE Tracking:** By mapping `cycle_time_ms`, we identify exactly where the 1% efficiency loss is happening.

5. Select Architecture Style

- **Selected Style:** Micro-kernel Architecture (Plug-in Style)

For the **On-Premises Factory Edge**, we will use a **Micro-kernel** style combined with **Layered Architecture**.

- **Core Logic:** The "Kernel" handles mission-critical safety tasks (*Safety Interlocks, PLC communication*). This stays lean to ensure **<5ms latency**.
- **Plug-ins:** The **Digital Twin** models for different machines (*Milling, Casting, Assembly*) are "Plug-ins." This allows us to update one machine's logic without stopping the whole factory.
- **Isolation:** If the "Acoustic Analysis" plug-in crashes, the core "Machine Stop" kernel remains active. This supports **99.99% Availability**.

6. Select Architecture Pattern

- **Selected Pattern:** Transactional Outbox Pattern

This pattern is critical for **Edge Autonomy** and **Zero Data Loss**.

How it works in ForgeSync

When a **Vibration Spike** or **Cycle Time** event occurs:

1. The system saves the event to a **Local Database (Outbox)** on the factory floor.
2. A separate "**Relay Service**" constantly **checks** this **Outbox** to **send data** to **Global HQ**.
3. **The WAN Outage Scenario:** If the **internet fails**, the **data** simply **stays** in the **Local Outbox**. When the **internet returns**, the **Relay Service** "**syncs**" everything in the **correct order**.

7. High-Level Design (HLD)

a. The 3-Layer ForgeSync HLD

The architecture is divided into three physical and logical zones:

- **Zone 1: The Shop Floor (On-Premises - Real Time)**
 - **Components:** Sensors (*Vibration, Acoustic*), PLCs, Edge Gateways.
 - **Flow:** *Raw signals → <5ms processing → Emergency Stop commands.*
- **Zone 2: The Factory Core (On-Premises - Local Control)**
 - **Components:** Local Server Cluster, Digital Twin Registry, Outbox Database.

- **Flow:** Event Aggregation → OEE Dashboarding → Local storage for 48-hour autonomy.
- **Zone 3: The Global Command Center (Corporate Cloud)**
 - **Components:** Global Analytics, AI Training, Financial Reporting.
 - **Flow:** Summarized KPIs → Long-term Predictive Maintenance (PdM) trends.

b. Communication Path

We use **MQTT** (*Message Queuing Telemetry Transport*) for **Zone 1** and **Zone 2** because it is **lightweight** and **handles** the "Firehose" of **1.2M events efficiently**.

We use **Secure gRPC** or **HTTPS** for **Zone 2** to **Zone 3** communication.

Note: **MQTT** (Message Queuing Telemetry Transport) is a lightweight, publish-subscribe messaging protocol designed for low-bandwidth, high-latency, or unreliable networks. It is widely used in IoT, real-time monitoring, and event-driven systems.

8. Low-Level Design (LLD)

a. Data Storage Design

To handle the **1.2M events/sec "Firehose"** we use a **Two-Tier Storage Strategy** at the **Factory Edge**:

- **Tier 1: Time-Series Database (TSDB)**
 - **Purpose:** Stores high-velocity sensor data (*Vibration, Acoustics*).
 - **Schema:** [timestamp, machine_id, sensor_type, value].
 - **Retention:** 7 days (*Raw data is deleted after 7 days to save local space*).
 - **Note:** A Time-Series Database (TSDB) is a database optimized to store, query, and analyze data points indexed by time.
- **Tier 2: Relational Database (RDBMS)**
 - **Purpose:** Stores the "Digital Twin" metadata (*Identity, Thresholds*) and the **Transactional Outbox**.
 - **Schema:** [machine_id, status, last_maintenance_date, oee_target].

b. Physical Deployment

- **Hardware: 3-Node High-Availability (HA) Server Cluster** located inside the Factory Server Room.
- **Storage:** NVMe SSDs to **support** the **high IOPS** (*Input/Output per second*) required for **1.2M events**.

9. Select Component, Feature, & Service

Component	Selected Technology	Core Purpose	The "Architect's Why" (Deep Dive)
Edge Gateway	Eclipse Kura	The Industrial Translator. It converts physical machine signals into software data.	PLCs speak "Industrial" (Modbus/OPC-UA), but our "Digital Twin" speaks "Cloud." This component bridges that gap in real-time .
Messaging Broker	EMQX (Enterprise MQTT)	The Digital Nervous System. It carries the 1.2M events/sec throughout the factory.	Traditional protocols like HTTP are too "heavy." MQTT is ultra-lightweight, allowing us to send 100 reports/sec from 12,000 sensors without crashing the network.
Stream Processor	Apache Flink (On-Prem)	The Intelligent Filter. It watches the data stream for "Grinding" or "Harmonic Spikes."	We cannot save 1.2M events/sec forever. Flink analyzes the data in flight , triggers the <5ms safety stop, and only saves the important parts.
High-Speed Storage (TSDB)	InfluxDB	The "Pulse" Recorder. It stores the raw, high-speed sensor readings (Hz, dB, °C).	A standard database (RDBMS) would lock up trying to write at this speed. InfluxDB is designed to handle "Time-Series" data—millions of points with timestamps.
Identity Storage (RDBMS)	PostgreSQL	The "History" Keeper. It stores machine IDs, safety thresholds, and the Outbox .	We need "ACID" compliance (100% data accuracy). If the power blinks, PostgreSQL ensures the Transactional Outbox doesn't lose the data we need for the HQ.
Orchestration	K3s (Lightweight K8s)	The Factory Brain Manager. It keeps all our software services running.	If our "Acoustic Monitor" service crashes, K3s restarts it automatically in seconds. This is how we guarantee 99.99% Availability .
Service Protocol	gRPC	The Instant Command. Used for machine-to-machine "Stop" signals.	gRPC is much faster than standard web calls. It allows the software to say "STOP" to the PLC in under 5ms , preventing motor melting.
Analytics Tool	Grafana	The Profit Dashboard. Displays the OEE health to the human operators.	If the manager can't see the 1% OEE gain , the project is a failure. Grafana turns the messy data in InfluxDB into a clear "Green/Red" dashboard.

Note: K3s is a lightweight, production-grade Kubernetes distribution created by Rancher (now part of SUSE), designed for resource-constrained environments like edge, IoT, and small on-prem setups.

Note: EMQX is a high-performance, distributed MQTT broker designed for massive-scale IoT and real-time messaging systems.

10. Architectural Decision Records (ADRs)

- **ADR-01: Choice of On-Premises over Cloud for Level 3**
 - **Decision:** All MES and **Digital Twin logic** will **reside** in the **factory**.
 - **Consequence:** Meets (**100% Autonomy**) and (**<5ms Latency**).
- **ADR-02: Choice of MQTT as Primary Protocol**
 - **Decision:** Use **MQTT instead of HTTP/REST** for **sensor data**.
 - **Consequence:** **Massive reduction** in **network overhead**; supports the "Firehose" scalability.
- **ADR-03: Use of Time-Series Database for Telemetry**
 - **Decision:** **Store sensor data** in **InfluxDB** rather than standard SQL.
 - **Consequence:** Allows for **high-speed analysis** of **Harmonic Spikes** and **Acoustic signatures**.

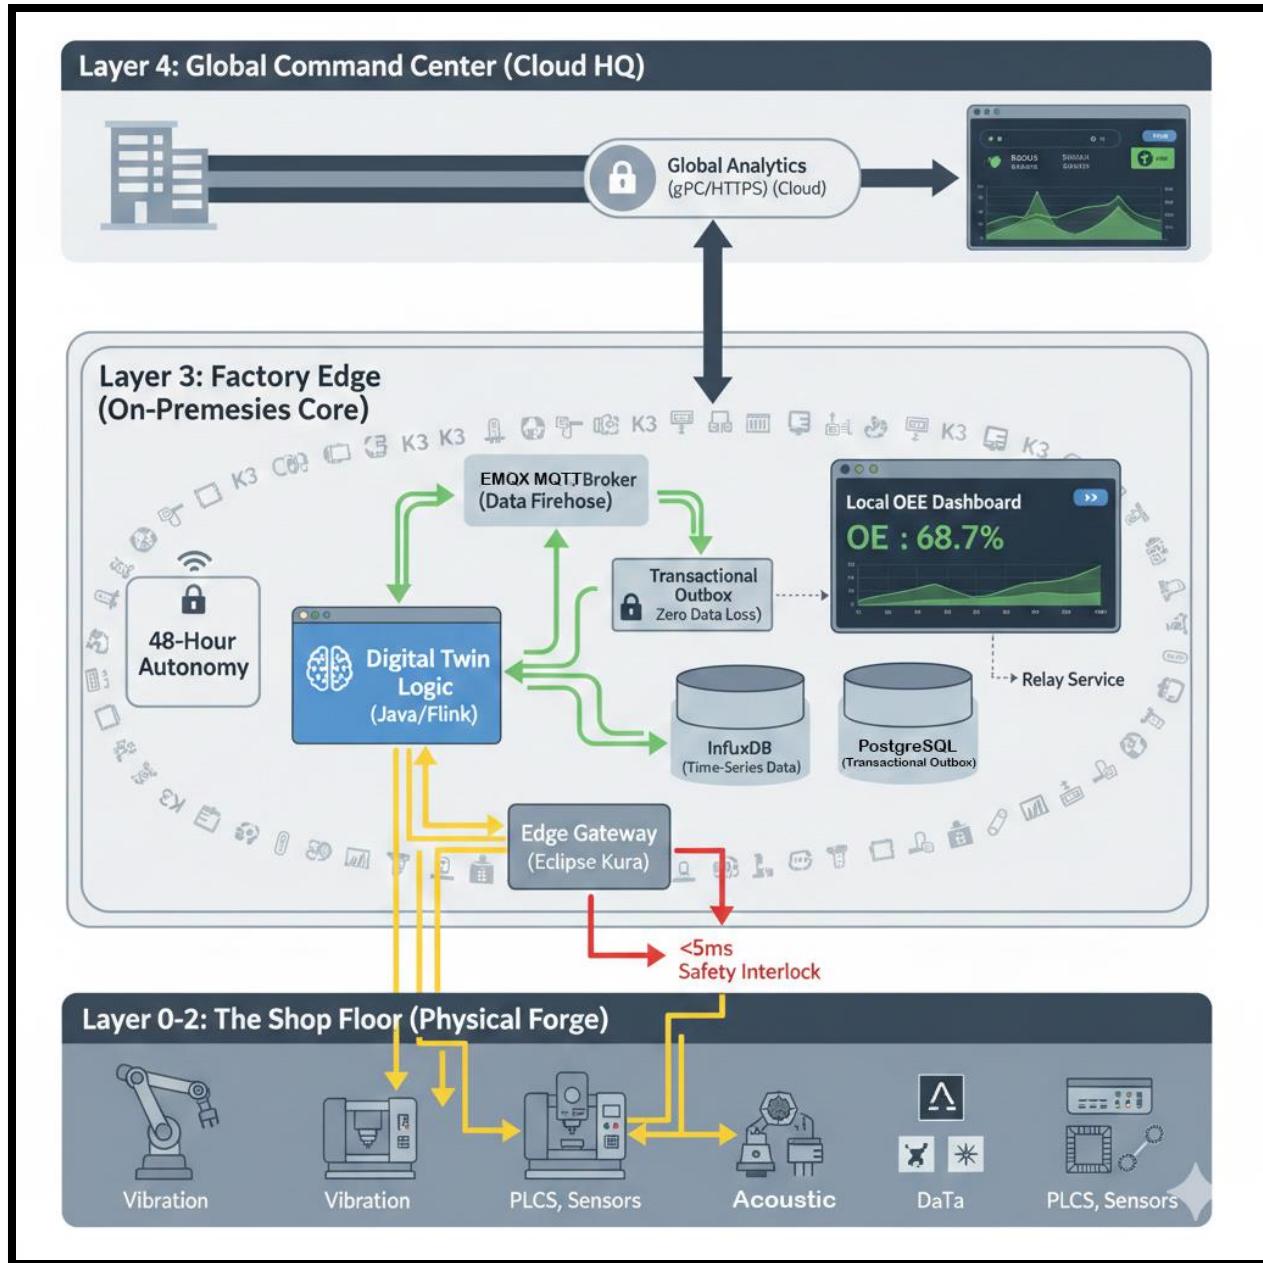
11. System Flow (The Vibration Spike)

- **Detection:** A physical **Vibration Sensor** **detects** a jump from 12Hz to **85Hz**.
- **Ingestion:** **Data** is **published** to the **On-Premises MQTT Broker** (*Latency: <1ms*).
- **Reaction:** The **Reactive Stream Processor** (*Flink*) identifies this as a "*Harmonic Spike*".
- **Action:** A **command** is **sent** via **gRPC** to the **PLC Interface** to **trigger** an **Emergency Stop**.
- **Synchronization:** The **event** is **written** to the **Transactional Outbox** (*PostgreSQL*).
- **Reporting:** Once the **WAN returns**, the **event** is **synced** to **HQ** for long-term AI training to prevent future spikes.

12.Final Architecture Picture

The Diagram should depict:

- **Bottom Layer:** The physical "Forge" with **Vibration/Acoustic sensors**.
- **Middle Layer (On-Premises):** The **MQTT Broker, Digital Twin Logic, and Local OEE Dashboard**. (The "Heart" of the system).
- **Side Layer:** The **Transactional Outbox** ensuring **data integrity**.
- **Top Layer (Global Cloud):** The **HQ receiving summarized OEE reports** via a secure tunnel.



1. Layer 0-2: The Shop Floor (The Physical Forge)

This is the bottom layer where the actual work happens.

- **The Senses:** These machines are equipped with **Vibration** and **Acoustic** sensors.
- **The Connection:** The yellow lines represent raw electrical signals and data being pulled from the machines and sent up to the "brain" (Layer 3).

2. Layer 3: Factory Edge (The On-Premises Core)

This is the "Heart" of your project. It is located physically inside the factory to ensure the system never stops.

- **Edge Gateway (Eclipse Kura):** This is the first stop. It translates machine language into digital data.
- **Safety Interlock (<5ms):** The **Red Arrow** pointing back down to the machines. This represents the ultra-fast emergency stop. If the Digital Twin hears "Grinding," it bypasses the cloud and stops the machine instantly.
- **The Firehose (EMQX MQTT Broker):** This handles the massive stream of 1.2M events per second.
- **Digital Twin Logic (Java/Flink):** This is the "Brain" where the virtual copy of the machine lives, evaluating health thresholds in real-time.
- **The Two Databases:**
 - **InfluxDB:** Storing the high-speed "Time-Series" waves (Vibration/Sound).
 - **PostgreSQL:** Storing the machine's "Identity" and the **Transactional Outbox**.
- **Local OEE Dashboard:** A screen showing **68.7% OEE**. This allows the local manager to see the \$45M improvement progress without needing an internet connection.
- **48-Hour Autonomy:** The "lock" icon on the left indicates that this entire gray box can run perfectly even if the internet cable is cut.

3. Layer 4: Global Command Center (Cloud HQ)

This is the top layer, representing the Corporate HQ.

- **The Secure Tunnel:** The thick gray pipes with a lock icon represent a secure connection (VPN/gRPC) that only sends **summarized** data.
- **Global Analytics:** HQ doesn't see every single vibration; they only see the "Health Stats" and the total OEE across all 4 factories.
- **ERP/AI Training:** This is where long-term trends are analyzed to improve the entire global fleet.

4. The Side Layer: The Transactional Outbox

Located in the middle-right, the **Transactional Outbox** acts as a "Data Safety Net."

- If the connection to the Top Layer (HQ) breaks, the data is safely queued here.
- The **Relay Service** (dotted line) waits patiently for the internet to return, then "pushes" the saved data up to HQ so that no production logs are ever lost.