

Capstone Project

Case Study

Project Overview: ShopSphere E-commerce Platform

Business Context

"**ShopSphere**," a rapidly expanding e-commerce retailer, specializes in personalized, high-value consumer goods (e.g., custom apparel, luxury electronics).

The company operates three geographically dispersed fulfillment centers (FCs) across different regions.

Due to regional data residency laws (e.g., GDPR, CCPA) and the sheer volume of global traffic, ShopSphere requires a decentralized architecture.

The core application logic and primary inventory systems must reside within regional data centers (DCs) to meet legal compliance and provide low-latency experiences to local customers.

The current legacy e-commerce system suffers from:

- **Data Silos:** Inventory data, user profiles, and order history are isolated in separate, inflexible databases.
- **Scalability Bottlenecks:** The monolithic architecture cannot handle Black Friday/holiday traffic spikes without crashing.
- **High Downtime Risk:** Deployments require full system restarts, leading to unacceptable downtime.

The Challenge: The "ShopSphere Core" Platform

You are tasked with designing the architecture for "**ShopSphere Core**," a new, decentralized platform that manages user traffic, product catalogs, and order processing.

The platform must operate primarily out of three separate Regional Data Centers (RDCs), each supported by an additional, compliant Disaster Recovery (DR) Data Center within its same legal jurisdiction and a Central Corporate Data Center (HQ-DC) for global reporting.

The core constraint is achieving near-global availability (99.95%) and low latency (**sub-100ms**) for all regional customers while ensuring data residency and localized compliance.

Business Scenario and Constraints

1. The Company and Industry

ShopSphere is a high-growth retailer in the personalized goods sector.

- **Core Products:** Custom fashion, premium gadgets, and exclusive digital subscriptions.
- **Key Driver:** Customer Experience and Conversion Rate. Every second of latency costs millions in lost sales.
- **Business Goal:** Achieve 99.95% Availability and increase Conversion Rate by 10% through high-speed, localized user experience.

2. Operational Environment (The Decentralization Constraint)

ShopSphere operates three primary Regional Pair Data Centers (*RDC-A, RDC-B, RDC-C*) and one centralized HQ-DC.

Region	Primary Data Center (RDC)	Compliant DR Data Center (DR-DC)	Jurisdiction
Region 1 (Europe)	RDC-A: Frankfurt, Germany	DR-DC-A: Dublin, Ireland	European Union
Region 2 (North America)	RDC-B: Ashburn, USA	DR-DC-B: Chicago, USA	United States
Region 3 (Asia-Pacific)	RDC-C: Singapore	DR-DC-C: Sydney, Australia	APAC Region
Global Reporting	HQ-DC: New York, USA		Central Management

- **Data Residency Mandate:** Core customer personal data, specific regional payment logs, and fulfillment data must reside within the respective RDC's physical boundaries (e.g., EU data must stay in RDC-A/DR-DC-A jurisdiction). Public cloud deployment is permissible only for CDN/Static Assets, but not for core transactional data.
- **Active-Passive DR Mandate:** The three DR-DCs are intended for Active-Passive compliant failover. They must be kept running and capable of taking over the primary RDC's transactional workload within 15 minutes to maintain the 99.95% availability target for that region.
- **RDCs (Regional Hubs):** Each RDC must operate independently (active-active deployment) to serve its local customers with 100ms latency. Each RDC has a fixed resource limit (e.g., 8 racks, 150kW power).
- **Global Network:** The WAN links between RDCs and the HQ-DC are high-speed but have inherent global latency (e.g., 200ms).

- **Existing Infrastructure:** The company uses Kubernetes for orchestration, standardizes on Linux OS, and has existing Global DNS and CDN providers that must be integrated.

3. Data Challenge

Data Source	Type	Volume/Velocity
Clickstream/Search Events	Small, structured JSON events	Extremely High Velocity (500,000 events/sec total)

4. Critical Technical Problems to Solve

Problem	Required Solution Implication
Data Residency & Low Latency	Core transactional data and application services must be geographically dispersed and operate independently.
High Availability (99.95%)	The platform must withstand a total failure of a primary RDC by rapidly failing over to its designated DR-DC.
High Traffic Spikes (Black Friday)	The platform must scale rapidly to handle 10x normal load using the existing Kubernetes cluster resources.
Transactional Integrity (Payments)	Financial records must be 100% accurate despite asynchronous communication between microservices.
Real-Time Personalization	Processing user behavior data must deliver personalized product recommendations within 50ms of page load.

3. The Assignment: 12 Detailed Tasks

Section 1: Strategic Mapping

- **Assignment 1: Business Vision to Technical Vision.**
- **Assignment 2: Functional & Non-Functional Requirements.**

Section 2: Architectural Selection

- **Assignment 3: Select Paradigm.**
- **Assignment 4: Select Model.**
- **Assignment 5: Select Architecture Style.**
- **Assignment 6: Select Architecture Pattern.**

Section 3: Technical Design & Flow

- **Assignment 7: High-Level Design (HLD).**
- **Assignment 8: Low-Level Design (LLD).**
- **Assignment 9: Component & Service Selection.**
- **Assignment 10: Create 3 ADRs (Architectural Decision Records).**

Section 4: Visualizing the Flow

- **Assignment 11: System Flow of Single Clickstream Event**
- **Assignment 11: Final Architecture Picture.**

Model Answer

1. Business Vision to Technical Vision

Business Vision	Technical Vision
Increase Conversion Rate by 10% via high-speed, localized user experience.	<ul style="list-style-type: none"> Implement a Global Server Load Balancing (GSLB) strategy combined with an Edge Caching Layer (CDN) to route users to the nearest Regional Data Center (RDC) and serve content with 100ms latency and system must handle 500,000 reads/second (IOPS) to manage traffic.
Achieve 99.95% Availability and recover from total RDC failure to prevent revenue loss during critical traffic spikes (e.g., Black Friday).	<ul style="list-style-type: none"> Design an Active-Passive Compliant Failover Strategy for each RDC/DR-DC pair, services must be fully restored (RTO) in <=15 minutes. This involves maintaining data replication of to the dedicated DR-DC within the same legal jurisdiction. Design a highly resilient, fault-tolerant deployment across all three RDCs, for rapid, resource-efficient horizontal scaling and isolated failure domains.
Ensure Data Residency and regulatory compliance (GDPR, CCPA) for customer data.	<ul style="list-style-type: none"> Implement a Geographically Partitioned Data Strategy where core customer and transactional data is strictly confined to and processed within its designated RDC and transactional data loss must be minimal, defined by an RPO of <=5 minutes, (due to inheritance WAN latency) and must guarantee 100% of sensitive records with no cross-region data replication for compliance.
Central Oversight Provide centralized global business reporting and analytics.	<ul style="list-style-type: none"> Anonymized Event Streaming. Establish a Central HQ-DC that consumes only anonymized and aggregated metrics, streamed asynchronously from all regions.

2. Functional and Non-Functional Requirements (FRs & NFRs)

Core Functional Requirements (FRs)

1. **Geo-Aware Routing & Catalog Display:** The platform must use **Global DNS (GSLB)** to route users to the **nearest RDC** and **dynamically display** the regional Product Catalog and **localized pricing** with a **latency** of **100ms**.
2. **Idempotent Order Transaction:** The **system** must process an **order transaction** with **ACID compliance** within the local RDC, ensuring **zero double-charges** or **lost orders**, and **commit** the **final financial record** before triggering downstream fulfillment processes.
3. **Real-Time Search & Recommendations:** The **platform** must **reliably ingest** and **process** user **Clickstream and Search Events** and **use** the resulting **personalized data** to **deliver product recommendations** on **core user pages** within **50ms** of **page load**.

Critical Non-Functional Requirements (NFRs)

1. **Availability (High Uptime & DR):** The platform must maintain **99.95% Availability** across all **core e-commerce services** (*Order, Catalog, User*). In the event of a **total primary RDC failure**, the system must guarantee a **Recovery Time Objective (RTO)** of **15 minutes** for that region via **automated failover** to its dedicated, **compliant DR-DC**.
2. **Data Residency (Compliance):** All customer **Order Transaction Data** and **personally identifiable information (PII)** must be **strictly and physically partitioned, processed, and stored only** within its **respective Regional Data Center (RDC/DR-DC)**, meeting the **mandatory regulatory requirements**.
3. **Scalability (Traffic Spikes):** The platform must **achieve horizontal scaling** of **critical services** (*Order and Catalog*) by **5x within 3 minutes** using the existing **Kubernetes cluster resources** in **each RDC** to **accommodate sudden, high-volume traffic events**.
4. **Latency (User Experience):** The **end-user HTTP response time** for **core transactional pages** (*e.g., Checkout, Product Detail Page*) must **not exceed 100ms** for **95% of requests (95)**, primarily **achieved through Edge and Service-level caching**.
5. **Data Integrity (RPO):** The **transactional data cluster** within **each RDC/DR-DC pair** must maintain a **Recovery Point Objective (RPO)** of **5 minutes** to minimize **financial data loss** during a compliant failover to the DR-DC.

3. Select Paradigm:

Paradigm Selection	Justification
Event-Driven Programming (EDP)	EDP is the most suitable mechanism and programming style for the high-volume, asynchronous , and decentralized nature of the ShopSphere platform.

4. Select Model:

a. Domain Model

The **Domain Model** for the ShopSphere E-commerce Platform is a conceptual blueprint that formally defines the **relationships, behaviors, and attributes** of its core business concepts, including:

Model Component	Example
Data Entities	<ul style="list-style-type: none"> User Order Product Inventory Payment Shipment
Relationships	<ul style="list-style-type: none"> User places Order (1: Many): Order contains Products (Many: Many) Order linked to Payment (1:1)
Attributes	<p>Order:</p> <ul style="list-style-type: none"> status total_amount shipping_address <p>Product:</p> <ul style="list-style-type: none"> price color description <p>User:</p> <ul style="list-style-type: none"> user_id Email residency_region
Behavior/ Processes	<ul style="list-style-type: none"> PlaceOrder() ReserveInventory() ProcessPayment() GenerateRecommendations() CompensateOrder()

b. Process Model

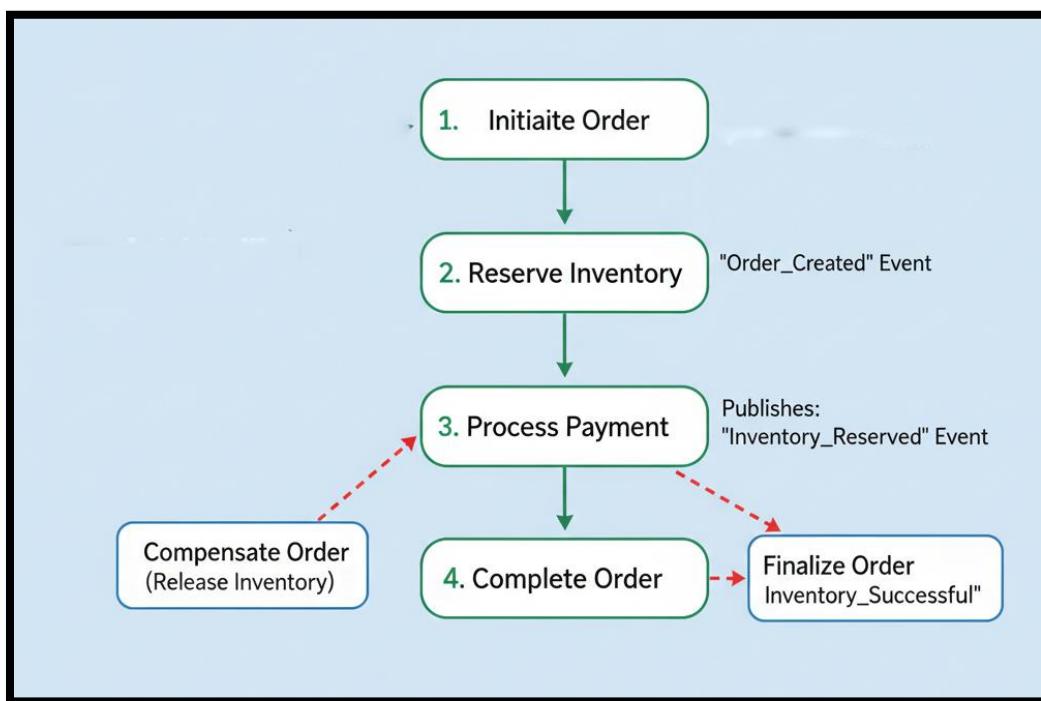
i. Order Flow

Step	Action	Path Type	Details & Business Logic
1. Initiate Order	Initiate Order	Synchronous	<ul style="list-style-type: none"> The transaction begins. This step typically involves the Order Service receiving the request, creating an initial record in its local database.
2. Reserve Inventory	Reserve Inventory	Synchronous / Event-Triggered	<ul style="list-style-type: none"> This step is triggered by the "Order_Created" Event. Catalog/Inventory Service consumes the event and attempts to logically deduct or reserve the necessary stock in its local database. This prevents overselling while the final payment is processed.
3. Process Payment	Process Payment	Synchronous	<ul style="list-style-type: none"> This is the most critical step. It involves the Order Service calling an external Payment Gateway to charge the customer. Upon successful inventory reservation, this step proceeds.
4. Complete Order	Complete Order	Synchronous	<ul style="list-style-type: none"> This step is triggered upon successful payment processing and inventory reservation. The Order Service updates its local record to a final status like PAID or COMPLETED.

ii. Failure and Success Paths

Path Type	Trigger Point	Action / Event Triggered	Service(s) Involved	Outcome and Rationale
Success Path	Successful Process Payment	Finalized Order (Inventory Successful)	Catalog/Inventory Service	<p>Final Commit:</p> <ul style="list-style-type: none"> The permanent deduction of inventory is performed. The Order Service updates the order status to COMPLETED. Transactional integrity is preserved successfully.

Compensation Path (Rollback)	Failure during Process Payment (e.g., card declined)	Compensate Order (Release Inventory)	Catalog/Inventory Service	Reversal/Consistency: <ul style="list-style-type: none"> This is the Compensating Transaction. The service releases the reserved inventory back to the available stock. This prevents inventory from being tied up by a failed order, maintaining data consistency.
------------------------------	---	--------------------------------------	---------------------------	---



5. Select Architectural Style:

Architecture Style Selection	Justification
Dominant Style: Microservices Architecture (MSA)	<ul style="list-style-type: none"> MSA is the mandatory style for the platform, providing the necessary boundaries for decentralization, scaling, and failure isolation.
Complementary Styles: Event-Driven Architecture and Layered Architecture	<ul style="list-style-type: none"> These styles are integrated to manage communication and provide internal structure, respectively.

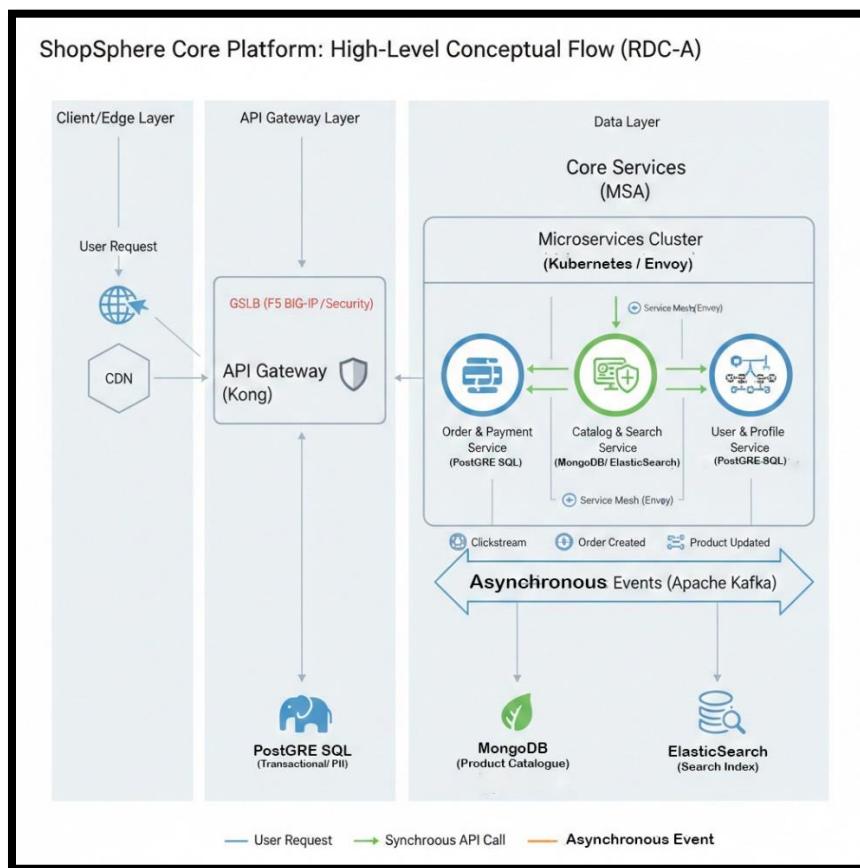
6. Select Architectural Pattern:

Pattern Selection	Justification
Saga Pattern (Choreography)	<ul style="list-style-type: none"> MANDATORY for Transactional Integrity. The Saga Pattern is essential to ensure ACID-like certainty for the Order Transaction across the many decoupled microservices.
Command Query Responsibility Segregation (CQRS)	<ul style="list-style-type: none"> MANDATORY for Performance and Scalability. CQRS separates the application's data models for read (query) operations from write (command) operations.
Circuit Breaker Pattern	<ul style="list-style-type: none"> CRITICAL for System Resilience and Availability. The Circuit Breaker wraps calls to external or unreliable services, preventing cascading failures.
Observer Pattern (via Message Broker)	<ul style="list-style-type: none"> CRITICAL for Real-Time State Notification. This behavioral pattern allows services to automatically receive updates when the state of an observed entity change.

7. High-Level Design (HLD): Conceptual Flow

The design is based on a standard **Event-Driven, Layered MSA**.

- Client/Edge Layer:** User requests enter via CDN → GSLB.
- API Gateway Layer:** Traffic is routed through the API Gateway (*Kong*) for policy enforcement (*rate limit*).
- Core Services Layer (MSA):** Requests are handled by the microservices (*Order*, *Catalog*, *User*) running on Kubernetes.
- Event/Communication Layer:** Services communicate asynchronously via Apache Kafka.
- Data Layer:** Polyglot Persistence (*PostgreSQL*, *MongoDB*, *Elasticsearch*) handles durability and optimized reads.



8. Low-Level Design (LLD) for Data Storage

a. PostgreSQL LLD (Order Transactions & PII)

PostgreSQL is chosen for its **ACID compliance** and mandatory **Data Residency** requirements.

i. Logical Design (Schema Details)

The **design** must **enforce data integrity** and efficiently **handle transactional workloads**.

Entity	Key Attributes	Data Type	Constraint	Rationale
Orders	order_id	UUID	Primary Key (PK)	Unique, immutable identifier
	user_id	UUID	Foreign Key (FK)	Links to the User table for compliance checks.
	status	ENUM	NOT NULL	Tracks state (CREATED, PAYMENT_PENDING, COMPLETED, COMPENSATED).
	total_amount	NUMERIC(10, 2)	NOT NULL	Financial integrity.

Users	user_id	UUID	Primary Key (PK)	Unique PII identifier.
	first_name	VARCHAR(50)	NOT NULL	Personal data for shipping/compliance.
	residency_region	VARCHAR(10)	NOT NULL	Mandatory field for Data Residency checks.
	created_at	TIMESTAMP	NOT NULL	Auditing.

ii. Physical Deployment (RDC-A and DR-DC-A)

The **deployment must ensure HA and compliant DR**.

Deployment Aspect	RDC-A (Active)	DR-DC-A (Passive)	Constraint Met
Cluster Topology	Primary Master + 2 Synchronous Replicas	One Standby Replica (Promotable)	for RDC failure.
Replication Mode	Synchronous Streaming (within RDC-A)	Asynchronous Streaming (RDC-A to DR-DC-A)	Zero Data Loss within RDC-A. RPO for DR, adhering to Data Residency .
Storage	High-performance SSD/NVMe attached volumes.	Standard SSD/EBS storage.	Low-latency I/O for high-volume transactions.
Failover	Automated failover managed by Patroni or similar cluster manager within RDC-A.	Manual/Semi-automated promotion of the standby replica.	High Availability (HA) .

9. Select Component, Feature, Service:

A. Critical Business-Logic Services (MSA Components per RDC/DR-DC)

These services are deployed via Kubernetes clusters in both the **RDC** (Active) and **DR-DC** (Passive).

Component/Service	Primary Purpose	Key Technologies Used	Key Features
User & Profile Service	Manages core customer PII, authentication/authorization, session state, and clickstream ingestion.	<ul style="list-style-type: none"> AuthN & AuthZ: Keycloak. RDBMS: PostgreSQL. Session: Redis. Messaging: Apache Kafka 	Authentication/Authorization: <ul style="list-style-type: none"> Implemented using OAuth 2.0 / OpenID Connect (OIDC). Manages user sign-up, login, and the issuance of JSON Web Tokens (JWTs). Profile Management: <ul style="list-style-type: none"> Ensures sensitive PII resides only within the local RDC/DR-DC

			<p>Jurisdiction using the clustered PostgreSQL.</p> <p>Session Management:</p> <ul style="list-style-type: none"> Application services (<i>e.g.</i>, User, Order, Catalog) do not save any user-specific data (session information) locally. All session data is stored externally in the centralized Redis cluster. Microservices fetch the required session details from Redis for every request. <p>Clickstream Publishing:</p> <ul style="list-style-type: none"> Captures raw user activity and publishes them directly to the local RDC's Kafka Cluster for high-throughput processing.
Catalog & Search Service	Provides low-latency product discovery and manages inventory visibility	<ul style="list-style-type: none"> Cache: Redis. Search: Elasticsearch. Database: MongoDB. 	<p>Localized Catalog Delivery:</p> <ul style="list-style-type: none"> Serves regional product data from the NoSQL database and Redis Cache (<i>100ms NFR</i>). <p>Search Index Management:</p> <ul style="list-style-type: none"> Maintains a dedicated, optimized Elasticsearch Index (<i>CQRS Read Model</i>) for ultra-fast query responses. <p>Flexible Schema for Products:</p> <ul style="list-style-type: none"> Utilizes the flexible schema of the NoSQL database to manage diverse product attributes and variants.
Order & Payment Service	Manages the transaction lifecycle, ensuring financial integrity and compliance	<ul style="list-style-type: none"> RDBMS: PostgreSQL Messaging: Apache Kafka 	<p>Idempotent Transaction Core:</p> <ul style="list-style-type: none"> Commits the financial record to the PostgreSQL (<i>local RDC</i>) with ACID compliance. <p>Saga Coordination:</p> <ul style="list-style-type: none"> Publishes and subscribes to events on the Kafka to manage the distributed Order Saga.

B. Critical Infrastructure/Data Components (Foundational Support)

Component/ Service	Role in Architecture	Key Technologies Used	Key Features and Constraint Fulfillment
Global Routing & CDN Layer	Manages ingress, traffic routing, and edge performance.	<ul style="list-style-type: none"> GSLB: Akamai Global Traffic Management (GTM) CDN: Cloudflare 	<p>Global Server Load Balancer (GSLB):</p> <ul style="list-style-type: none"> Directs users to the nearest, healthiest RDC based on geo-proximity. <p>CDN:</p> <ul style="list-style-type: none"> Caches static assets at the edge, supporting the 100ms latency NFR.
Data Persistence Cluster (RDC-DR Compliant)	Provides storage and implements the Active-Passive DR strategy .	<ul style="list-style-type: none"> RDBMS Replication: PostgreSQL Streaming Replication (<i>support both Synchronous and Asynchronous replication</i>) NoSQL: Redis Cluster 	<p>Asynchronous DR Replication:</p> <ul style="list-style-type: none"> Uses PostgreSQL Streaming Replication to copy transactional logs from RDC → DR-DC for compliant failover (RPO 5min). <p>RDBMS Cluster (HA):</p> <ul style="list-style-type: none"> Uses synchronous replication within the primary RDC for instant HA.
Local Load Balancing & API Gateway	Manages internal traffic distribution, security, and external API exposure within each RDC.	<ul style="list-style-type: none"> Load Balancer: F5 Load Balancer API Gateway: Kong 	<p>L7 Load Balancer:</p> <ul style="list-style-type: none"> Distributes traffic to the API Gateway instances. <p>API Gateway:</p> <ul style="list-style-type: none"> Enforce Authentication, Rate Limiting, and Throttling (e.g., using <i>Redis for counters</i>) to protect internal services during traffic spikes.
Message Broker Cluster	Provides the asynchronous communication backbone.	<ul style="list-style-type: none"> Messaging: Apache Kafka 	<p>Durable Messaging:</p> <ul style="list-style-type: none"> Guarantees message delivery, enabling the EDA/Saga pattern. <p>High Throughput:</p> <ul style="list-style-type: none"> Handles high-velocity Clickstream data ingestion without failure.

C. Critical Central HQ-DC Components (Global Management & Reporting)

Component/ Service	Role in Architecture	Key Technologies Used	Key Features and Constraint Fulfillment
Global Observability Hub	Centralized tooling for technical operations and management.	<ul style="list-style-type: none"> Monitoring: Prometheus & Grafana Logging: ELK Stack (Elasticsearch, Logstash, Kibana) 	Global Health Dashboard: <ul style="list-style-type: none"> Provides a single, unified view of the system's operational health, latency, and resource utilization across all RDCs and DR-DCs.

10. Architectural Decision Records (ADR):

a. ADR 001: Adopt CQRS for Catalog & Search Service

Decision	The Catalog & Search service will utilize CQRS, separating transactional (MongoDB Write Model) from search (Elasticsearch Read Model) data stores.
Status	Accepted
Context	Need to achieve 100ms query latency and high throughput while maintaining complex product data integrity.
Consequences	<ul style="list-style-type: none"> Pros: Optimizes the read path for speed and complex queries; search engine failure does not impact transactional writes. Cons: Introduces eventual consistency between the Write and Read models, managed via a dedicated Kafka pipeline.

b. ADR 002: Adopt PostgreSQL as the Transactional System of Record

Decision	PostgreSQL is selected as the mandatory database for the Order & Payment Service and PII storage within the User & Profile Service.
Status	Accepted
Context	<ul style="list-style-type: none"> Financial records and PII require strict durability and transactionality (ACID compliance). Furthermore, robust, open-source Streaming Replication is required for compliant DR.
Consequences	<ul style="list-style-type: none"> Pros: Guaranteed transactional integrity and consistency. Robust, proven mechanism for compliant RPO 5min replication to the DR-DC. Cons: Less flexible schema than MongoDB; requires significant operational expertise for maintenance and tuning at high scale.

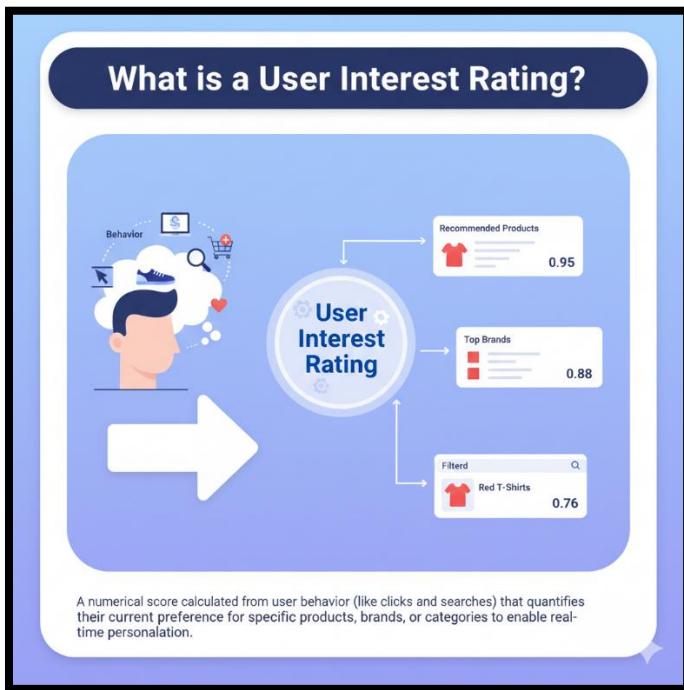
c. ADR 003: Adopt MongoDB as the Primary Catalog Write Store

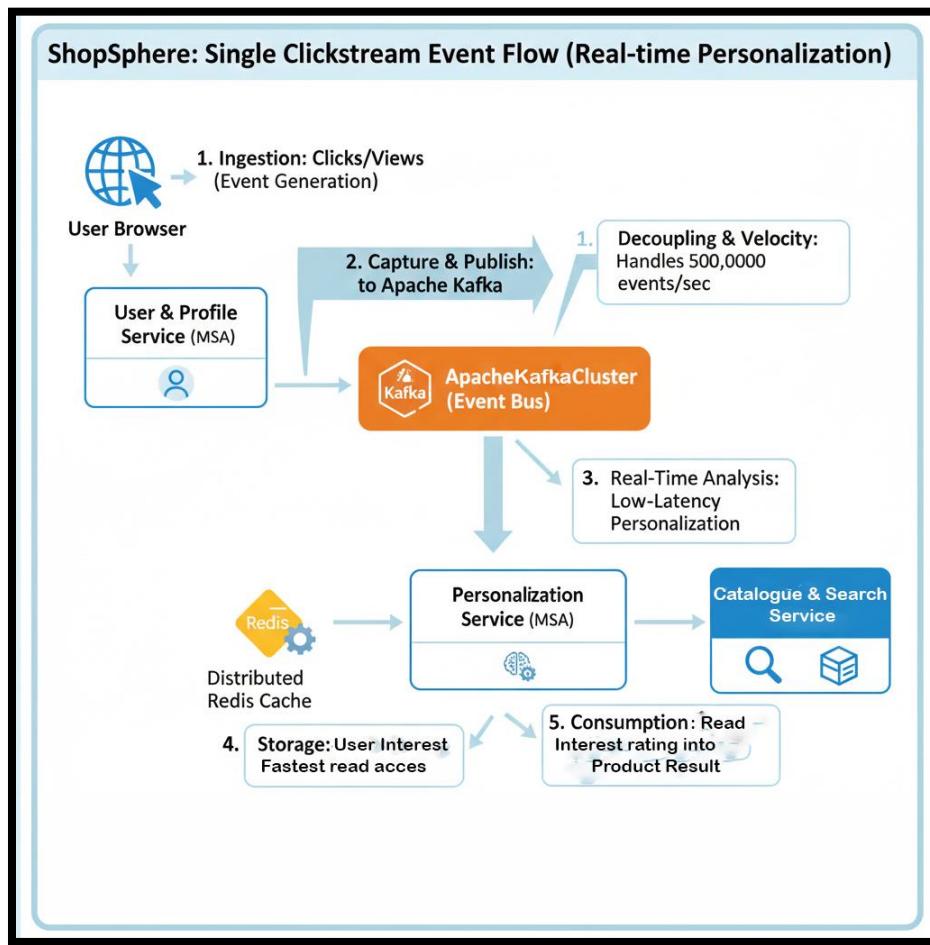
Decision	MongoDB is selected as the primary database for the Catalog Service's Write Model, leveraging the CQRS pattern.
Status	Accepted
Context	The Product Catalog involves complex, highly varied, and frequently changing product attributes which do not fit a rigid relational schema.
Consequences	<ul style="list-style-type: none"> Pros: Provides necessary flexible schema for complex e-commerce product variants. High horizontal scalability for the Catalog's data volume. Cons: Eventual consistency between MongoDB (Write) and Elasticsearch (Read) must be carefully managed via Kafka. Cannot be used for ACID financial records.

11. System Flow: Single Clickstream Event

Step	Component	Action/Technology	Rationale and Constraint
1. Ingestion	User Browser	Clicks/Views (Event Generation): The user interacts with the application (<i>e.g., product view, search action</i>)	Generates the raw event data at the origin
2. Capture & Publish	User & Profile Service	Ingests event → Immediately publishes raw event to Apache Kafka via a high-throughput producer API.	Decoupling and Velocity: Kafka handles 500,000 events/sec velocity.
3. Processing	Personalization Service (MSA)	Consumes the event from Kafka → Applies stream processing rules (<i>e.g., aggregates events, calculates user interest rating</i>).	Real-Time Analysis: Uses Event Driven Programming to process data instantly for low-latency personalization.
4. Storage	Personalization Service	Stores the updated user interest rating in the Distributed Redis Cache .	Speed: Redis provides the fastest possible read access for the recommendation engine .
5. Consumption	Catalog & Search Service	Reads the interest rating from the Redis Cache during a search/view request.	Low Latency (100ms): Uses the in-memory cache to retrieve personalization data

		instantly to integrate into the product results.
--	--	--





12.Final Architecture Picture: Decentralized Multi-Region

The final architecture synthesizes all components and communication strategies, demonstrating **Data Residency**, **Scalability**, and **Resilience** across the three operational sites (RDC, DR-DC) and the central hub (HQ-DC).

