

## Case Study

### "Bharat-Krishi Connect": National Precision Agriculture & Unified Data Platform

#### 1. Detailed Business Scenario

**The Organization:** The National Agricultural Research & Development Board (NARDB).

**The Mission:** To modernize India's agricultural backbone by launching "Bharat-Krishi Connect." This platform aims to provide **Real-time Crop Advisory** and **Automated Resource Management** (*water/pesticides*) to **100 million farmers across the country**.

**The Mission:** Build a unified, multi-tenant platform that serves **140 million landholdings**.

#### The Phased Roadmap:

The platform is designed to scale horizontally as more states and landholdings are onboarded.

- **Phase 1 (Pilot): Implementation** across **4 Southern States (Tamil Nadu, Kerala, Karnataka, and Telangana)**.
  - **Target:** **25 million Farmers** and **30 million Landholdings**.
  - **Focus:** Managing diverse topography from **coastal belts** to the **Deccan Plateau**.
- **Phase 2 (Scale):** Expansion to **100 Million Farmers** (*reaching approximately 50% of the national footprint*).
- **Phase 3 (Full Rollout):** Final **unified nation**.

The **platform** must move from "**Manual Monitoring**" to "**Autonomous Farming**".

#### Business Goals & Key Results (KRs):

- **Yield Increase:** **Improve national crop productivity** by **18%** through AI-driven **sowing recommendations**.
- **Water Conservation:** **Reduce** agricultural **water consumption** by **25%** using **automated IoT-gated irrigation**. **Success** requires a **<500ms end-to-end "Sensor-to-Valve"** control loop.
  - **Core Mechanism:** Rain-Preemption Logic. The **system** must automatically issue a "**SKIP\_IRRIGATION**" **command** to all relevant **farm gateways** if the **localized IMD (India Meteorological Department) rain probability exceeds 70%**.

- **Market Access:** Provide **real-time "Mandi" price** transparency to farmers within **2 seconds** of **price updates**.
- **Disaster Mitigation:** Reduce crop loss from **pests/weather** by **sending alerts** to **2.5 million+ concurrent users** within **60 seconds** of an event detection.

### The Technical Challenge: Extreme Scale & Resilience

- **Astronomical Scale:** **Managing 150+ million hectares** requires a system that can **ingest 2 million+ messages per second** at **peak** national rollout. The **storage** layer must **handle petabytes** of **time-series data** while remaining searchable for historical yield trends.
- **Rural Connectivity Gaps:** Despite 5G growth, the **last mile** in rural India remains **unstable**. The **architecture** must **solve** the **Intermittent Connectivity**.
- **Real-Time Performance:** "**Actionable Insights**" (*e.g., detecting a locust swarm or a pump failure*) **lose value** if **delayed**.
- **Data Integrity (The RPO Challenge):** The system must survive regional disasters without losing critical "State" data. An **RPO of < 5 seconds** is **mandatory** to ensure that the **"Last Known State"** of **30M+** valves is **consistent** across **primary** and **backup sites**.

## 2. Operational & Technical Constraints

### A. Data Residency (The Legal Constraint)

- **Sovereign Boundary:** All **data** must be **hosted** on **servers** physically located **within India**. No data or metadata can be processed in overseas regions (*e.g., US-East or EU-West*).
- **Placement Flexibility:** **Data can** be **moved freely** between any Indian data centers (*e.g., Mumbai, Chennai, Delhi, or Hyderabad*) for load balancing or disaster recovery.

### B. High Availability & Disaster Recovery

- **99.99% Uptime:** Agriculture is **time-sensitive**. A **system outage** during the **sowing season** could lead to **national food shortages**.
- **15-Minute RTO (Recovery Time Objective):** **Maximum time allowed** to **switch operations** from the Primary Hub to the DR site during a total outage.
- **< 5-Second RPO (Recovery Point Objective):** The system must guarantee that **no more than 5 seconds of sensor data or command logs are lost** during a failover. This prevents "Ghost Watering" (double-irrigation) caused by data gaps during a site switch.

- **Active-Active/Passive:** The system **must survive** a **total outage** of a **primary region** (e.g., Mumbai) by **failing over** to a **secondary region** (e.g., Hyderabad) within **15 minutes**.

### C. Technical Friction Points

- **Legacy Integration:** The **platform** must **pull data** from **existing state-level weather stations** that use older SOAP APIs.
- **Peak Loads:** During **monsoon** or **harvest periods**, **traffic spikes 20x compared to the off-season**.

## 3. The Assignment: 12 Detailed Tasks

Students must act as the Lead Solution Architect to complete the following:

### Section 1: Strategic Mapping

- **Assignment 1: Business Vision to Technical Vision.**
- **Assignment 2: Functional & Non-Functional Requirements.**

### Section 2: Architectural Selection

- **Assignment 3: Select Paradigm.**
- **Assignment 4: Select Model.**
- **Assignment 5: Select Architecture Style.**
- **Assignment 6: Select Architecture Pattern.**

### Section 3: Technical Design & Flow

- **Assignment 7: High-Level Design (HLD).**
- **Assignment 8: Low-Level Design (LLD).**
- **Assignment 9: Component & Service Selection.**
- **Assignment 10: Create 3 ADRs (Architectural Decision Records).**

### Section 4: Visualizing the Flow

- **Assignment 11: Create System Flow.**
- **Assignment 12: Final Architecture Picture.**

## Model Answer

### 1. Mapping Business Goals to Technical Pillars

This section translates the high-level KR (*Key Results*) into structural pillars that will guide the design of the **Bharat-Krishi Connect** platform.

Business Goal	Technical Pillar	Success Metrics & KPIs
National Productivity Growth	Unified Analytical Intelligence	<ul style="list-style-type: none"> <li><b>Data Consolidation:</b> Integrate telemetry from <b>30M landholdings</b> (<i>Phase 1</i>) into a <b>standardized schema</b> for cross-regional ML modeling.</li> <li><b>Yield Optimization:</b> Deliver precise sowing windows to <b>achieve an 18% productivity increase</b> by synthesizing soil, drone, and historical datasets.</li> </ul>
National Water Stewardship	Closed-Loop Autonomous Control	<ul style="list-style-type: none"> <li><b>Precision Actuation:</b> Maintain <b>&lt; 500ms end-to-end latency</b> for sensor-to-valve triggers.</li> <li><b>Conservation Target:</b> Drive a <b>25% reduction</b> in <b>water usage</b> by transitioning from manual/scheduled timers to real-time, sensor-driven automation.</li> </ul>
Farmer Market Empowerment	Real-Time Information Symmetry	<ul style="list-style-type: none"> <li><b>Low-Latency Propagation:</b> Utilize high-speed messaging patterns to <b>broadcast</b> national Mandi <b>price changes</b> in <b>&lt; 2 seconds</b>.</li> </ul>
Emergency Risk Mitigation	High-Velocity Elastic Messaging	<ul style="list-style-type: none"> <li><b>Concurrent Notification:</b> Scale infrastructure to <b>broadcast critical pest</b> or <b>weather alerts</b> to <b>2.5M+ concurrent users</b> within <b>60 seconds</b> of detection.</li> </ul>
National Mission Criticality	Centralized Regional Resiliency	<ul style="list-style-type: none"> <li><b>Continuous Availability:</b> Maintain <b>99.99% system uptime</b> through primary hub and local DR pairing (e.g., <i>Chennai → Madurai</i>).</li> <li><b>Disaster Recovery:</b> Commit to a <b>15-minute RTO</b> (<i>Recovery Time</i>) and <b>&lt; 5s RPO</b> (<i>Data Loss window</i>) during regional outages.</li> </ul>
Rural Infrastructure Resilience	Decentralized Edge Autonomy	<ul style="list-style-type: none"> <li><b>Offline Continuity:</b> Empower farm gateways to maintain <b>100% irrigation functionality</b> for up to <b>48 hours</b> during a <b>total cloud-link blackout</b>.</li> </ul>
Massive National Scale	Hyper-Scale Horizontal Elasticity	<ul style="list-style-type: none"> <li><b>Linear Scalability:</b> Architect a <b>system</b> capable of onboarding <b>25M to 140M landholdings</b> and millions of IoT devices without latency degradation or data loss.</li> </ul>

## 2. Functional (FR) & Non-Functional Requirements (NFR)

### I. Functional Requirements (FR)

Functional requirements define **what** the **system** must **do** to support the **140 million farmers** and the **water conservation mission**.

ID	Requirement Category	Detailed Description
FR-1	Autonomous Irrigation	The <b>system</b> must trigger " <b>Open/Close</b> " commands to field valves based on <b>soil moisture</b> thresholds and <b>crop-specific hydration plans</b> .
FR-2	Predictive Preemption	The <b>system</b> must ingest weather forecast data and automatically abort scheduled irrigation if rain probability exceeds <b>70%</b> in a <b>6-hour window</b> .
FR-3	Market Transparency	The <b>platform</b> must <b>push "Mandi" price</b> updates to the <b>mobile app</b> for crops registered in the <b>farmer's profile</b> .
FR-4	Legacy Integration	The <b>system must pull data</b> from <b>state-level stations</b> via <b>SOAP/XML</b> and normalize it into a <b>unified national schema</b> .
FR-5	Emergency Alerting	The <b>system</b> must <b>detect pest outbreaks</b> via <b>drone imagery analysis</b> and <b>broadcast regional alerts</b> to all <b>farmers</b> in the <b>affected "Grid Cell"</b> .

### II. Non-Functional Requirements (NFR)

Non-functional requirements define **how** the system performs. These are the "Design Constraints" that ensure the **25% water saving** and **99.99% uptime**.

#### A. Performance & Scale (The "Krishi" Load)

- **NFR-1 (Latency):** The **end-to-end "Sensor-to-Valve"** control loop (*Actuation Latency*) must be **less than 500ms** to prevent over-watering or pipe bursts.
  - **The "Latency Budget" Breakdown**
    - To **achieve** the **< 500ms** target, the technical requirements for sub-components are as follows:

Component	Time	Technical Reason (The "Why")
Ingestion (Up)	150ms	<b>Signal Travel:</b> Covers the time for a sensor in a rural field to reach the nearest cell tower and travel via fiber to the Regional Hub.
Logic (Brain)	100ms	<b>Decision Making:</b> The time for the server to check the "Rain Rule," verify the farmer's ID, and decide whether to turn the water on or off.
Downlink (Down)	150ms	<b>Return Trip:</b> The time for the "Stop" command to travel back from the Hub through the 4G/NB-IoT network to the farm gate.

<b>Actuation (Physical)</b>	<b>100ms</b>	<b>Mechanical Move:</b> The physical time a relay switch or solenoid valve takes to physically flip from an "Open" to "Closed" position.
<b>Total</b>	<b>500ms</b>	<b>The Safety Limit:</b> The maximum delay allowed to ensure precision irrigation and prevent pipe damage (Water Hammer).

- **NFR-2 (Throughput):** The ingestion layer must support a sustained load of **150,000 messages per second (MPS)** and an **elastic peak** of **500,000 MPS** during monsoon surges.
  - **The Capacity Calculation (The "How Many")**
    - The **150,000 MPS** is the "Sustained Baseline" **required to keep** the **data fresh** enough for the **logic engine** to make **smart decisions**.
    - **Population: 30 Million Landholdings** (Phase 1).
    - **Freshness Goal:** We need an **update** every **200 seconds** (3.3 mins). The **200-second interval** tells the Hub "**What is happening**" often enough, so that when a command is triggered, it can be executed in **< 500ms**.
    - **The Math:  $30,000,000 \text{ Devices} \setminus 200 \text{ Seconds} = 150,000 \text{ Messages Per Second (MPS)}$**
- **NFR-3 (Price Propagation):** Market price updates must be **fanned out** to **2 million concurrent users** in **less than 2 seconds**.

## B. Availability & Resilience

- **NFR-4 (Uptime):** The **system** must achieve **99.99% availability** (*max 52 minutes of downtime per year*).
- **NFR-5 (Disaster Recovery):** In the event of a **regional failure** (e.g., *Mumbai DC*), the system must **failover** to the **secondary region** (e.g., *Hyderabad*) within an **RTO of 15 minutes**.
- **NFR-6 (Data Integrity):** The **Recovery Point Objective (RPO)** for **water consumption logs** must be **less than 5 seconds** to ensure accurate national water accounting.

## C. Sovereignty & Security

- **NFR-7 (Data Residency):** All **PII**, geospatial data, and moisture logs must be physically **stored** within the **Indian Sovereign Boundary**.
- **NFR-8 (Edge Autonomy):** The **local farm gateway** must **retain** enough "**State**" to perform **autonomous irrigation** for **48 hours** during a total cloud-connectivity blackout.



[www.ahmadzahoory.com](http://www.ahmadzahoory.com)

Ahmad Majeed Zahoory

### 3. Architectural Paradigm Selection

For the "Bharat-Krishi Connect" platform, we select the **Functional Reactive Paradigm**, where **Reactive** as the "**Logistics**" (*how we move data*) and **Functional** as the "**Decision**" (*how we process data*).

Paradigm	Role in the Platform	Simple Reason (The "Why")	Impact on NFRs
Reactive	The High-Speed Pipeline	<b>Non-Stop Movement:</b> It ensures the server never "waits". It can handle <b>150,000 messages</b> at once without crashing or getting stuck.	<b>Solves Scale:</b> Handles the 5M MPS peak and 2M concurrent users.
Functional	The Instant Decision	<b>Stateless Math:</b> It treats the " <b>70% Rain Rule</b> " as a <b>fast math equation</b> . Since <b>it doesn't have</b> to "look up" or " <b>change</b> " <b>data</b> , it is <b>incredibly fast</b> .	<b>Solves Latency:</b> Keeps the logic step under the <b>100ms budget</b> .

#### The "Winning Combo" Effect

- **Reactive** handles the **Quantity**: It manages the **massive crowd of 30 million sensors** trying to talk at the same time.
- **Functional** handles the **Quality**: It ensures that every decision (*like "Skip Irrigation"*) is **calculated** with **100% accuracy** in a **fraction of a second**.

### 4. Select Model

#### a) Domain Model (Logical Entities)

This model defines the core objects required to manage **30 million landholdings** and **25 million farmers**.

Entity	Core Attributes	Relationship
<b>Farmer</b>	CitizenID, Name, Contact, Language Preference	Owns one or more <b>Landholdings</b> .
<b>Landholding</b>	LandUID, Geo-Polygon (GIS), Soil Type, Primary Crop	Host for one <b>Gateway</b> and multiple <b>Sensors</b> .
<b>Gateway</b>	DeviceID, Firmware Version, Connectivity Status	Acts as the communication bridge for a <b>Landholding</b> .
<b>Telemetry</b>	ReadingID,	Linked to a specific <b>Gateway</b> or <b>Sensor</b> .

	Timestamp, Moisture, Temperature, pH	
<b>Commodity</b>	Name (e.g., Rice), Variety, Grade	Subject of <b>Mandi Price</b> updates.
<b>Mandi Hub</b>	HubID, Location Name, District	Source of real-time price fluctuations.
<b>Weather Grid</b>	GridID, Boundary Coordinates, Rain Probability	Overlays across multiple <b>Landholdings</b> .

## b) Process Model

### i. Autonomous Actuation Loop

**Target:** < 500ms (Ensuring Precision Irrigation)

Step	Action	Logic / Requirement
<b>Ingestion</b>	Sensor Data Uplink	Farm Gateway pushes current moisture status to the Regional Hub.
<b>Contextualize</b>	Fetch Environment State	System identifies the <b>Weather Grid</b> associated with the Landholding.
<b>Evaluate</b>	Apply "Rain Rule"	Logic Engine checks: <i>Is Rain Prob &lt; 70% AND Moisture &lt; 20%</i> ?
<b>Command</b>	Issue Downlink	If condition is met, Hub generates an "Open Valve" instruction.
<b>Execution</b>	Physical Relay Flip	Gateway receives command and triggers the mechanical water valve.

### ii. Market Access (Price Transparency)

**Target:** < 2 Seconds (Market-to-Mobile Delivery)

Step	Action	Logic / Requirement
<b>Capture</b>	Price Ingestion	Central eMandi system receives a price update for a specific <b>Commodity</b> .
<b>Standardize</b>	Schema Normalization	Data is converted into a standard unit (e.g., Price per Quintal in INR).
<b>Targeting</b>	User Segmentation	System finds all <b>Farmers</b> interested in that Commodity or region.
<b>Broadcast</b>	Reactive Push	Real-time push transmission to the active mobile app sessions.

<b>Notify</b>	Device Presentation	The farmer's app displays the updated price and trend analysis.
---------------	---------------------	---

### iii. Disaster Mitigation (Emergency Alerting)

**Target: < 60 Seconds** (Event-to-Masses Notification)

Step	Action	Logic / Requirement
<b>Detect</b>	Anomaly Identification	Drone AI or Weather Sensor detects a disaster event (e.g., Pest Swarm).
<b>Scope</b>	Geospatial Mapping	System identifies all <b>Landholdings</b> within the danger polygon.
<b>Prioritize</b>	Emergency Routing	Alert enters a high-priority queue, bypassing all routine sensor traffic.
<b>Fan-out</b>	Mass Notification	Simultaneous broadcast to <b>2.5M+ concurrent users</b> in the affected zone.
<b>Action</b>	Mitigation Guidance	Notification includes localized steps to minimize crop damage.

## 5. Select Architectural Style

The **Architectural Style** defines the "macro" view of the system.

We use **Microservices** to divide the project into independent, autonomous units that communicate over a network.

Component	Responsibility	Phase 1 Real-World Example	Rationale for this Style
<b>Control Service</b>	Manages the <b>500ms</b> sensor-to-valve loop.	A sensor in a <b>Telangana</b> cotton farm detects dry soil and triggers a pump.	<b>Resilience:</b> If the price service crashes, the water system stays alive.
<b>Market Service</b>	Manages the <b>2-second</b> price fan-out.	Pushing a sudden "Turmeric" price surge to 15 million mobile apps.	<b>Elasticity:</b> We can scale this service to 100 nodes during harvest peaks.
<b>Alert Service</b>	Manages the <b>60-second</b> emergency broadcast.	Detecting a locust swarm in <b>North Karnataka</b> and alerting nearby farms.	<b>Priority:</b> Emergency traffic is handled separately from routine soil data.
<b>Identity Service</b>	Manages <b>Aadhaar</b> and Land GIS records.	Verifying that a specific farmer owns a specific plot before acting.	<b>Security:</b> Isolates sensitive PII (Sensitive Data) from high-velocity IoT traffic.

## 6. Select Architectural Pattern

### a) Hexagonal (Ports & Adapters)

- **Best for:** The **Control Service** (*Sensor-to-Valve logic*).
- This **pattern** is used where **Latency (< 500ms)** and **Logic Protection** are the highest priorities.

Pattern Layer	Role in Bharat-Krishi	Real-World Example (Control Service)	Rationale
<b>Domain (The Core)</b>	The "Pure" Business Logic.	The code calculating the <b>"70% Rain Rule"</b> using only moisture and forecast variables.	<b>Speed:</b> Runs in RAM with zero external dependencies. Keeps the logic step under <b>100ms</b> .
<b>The Ports</b>	Contractual Interfaces.	A standard "socket" for receiving soil data or sending "Stop Pump" commands.	<b>Independence:</b> The core logic doesn't care if the sensor is 4G, 5G, or Satellite.
<b>Input Adapters</b>	Technology-specific entry.	An <b>MQTT Adapter</b> that translates raw sensor signals from the farm into a "Moisture Event."	<b>Flexibility:</b> Allows the same core logic to work with different hardware types across Phase 1.
<b>Output Adapters</b>	Technology-specific exit.	A <b>Redis Adapter</b> that fetches the latest weather status from a fast cache instead of a slow DB.	<b>Latency Shield:</b> Prevents a slow database from "blocking" the decision to turn off a valve.

### b) CQRS (Command Query Responsibility Segregation)

- **Best for:** The **Market Service** (*Mandi Price updates*).
- This **pattern** is used where **Read Volume (15M users)** is **vastly different** from **Write Volume (Mandi updates)**.

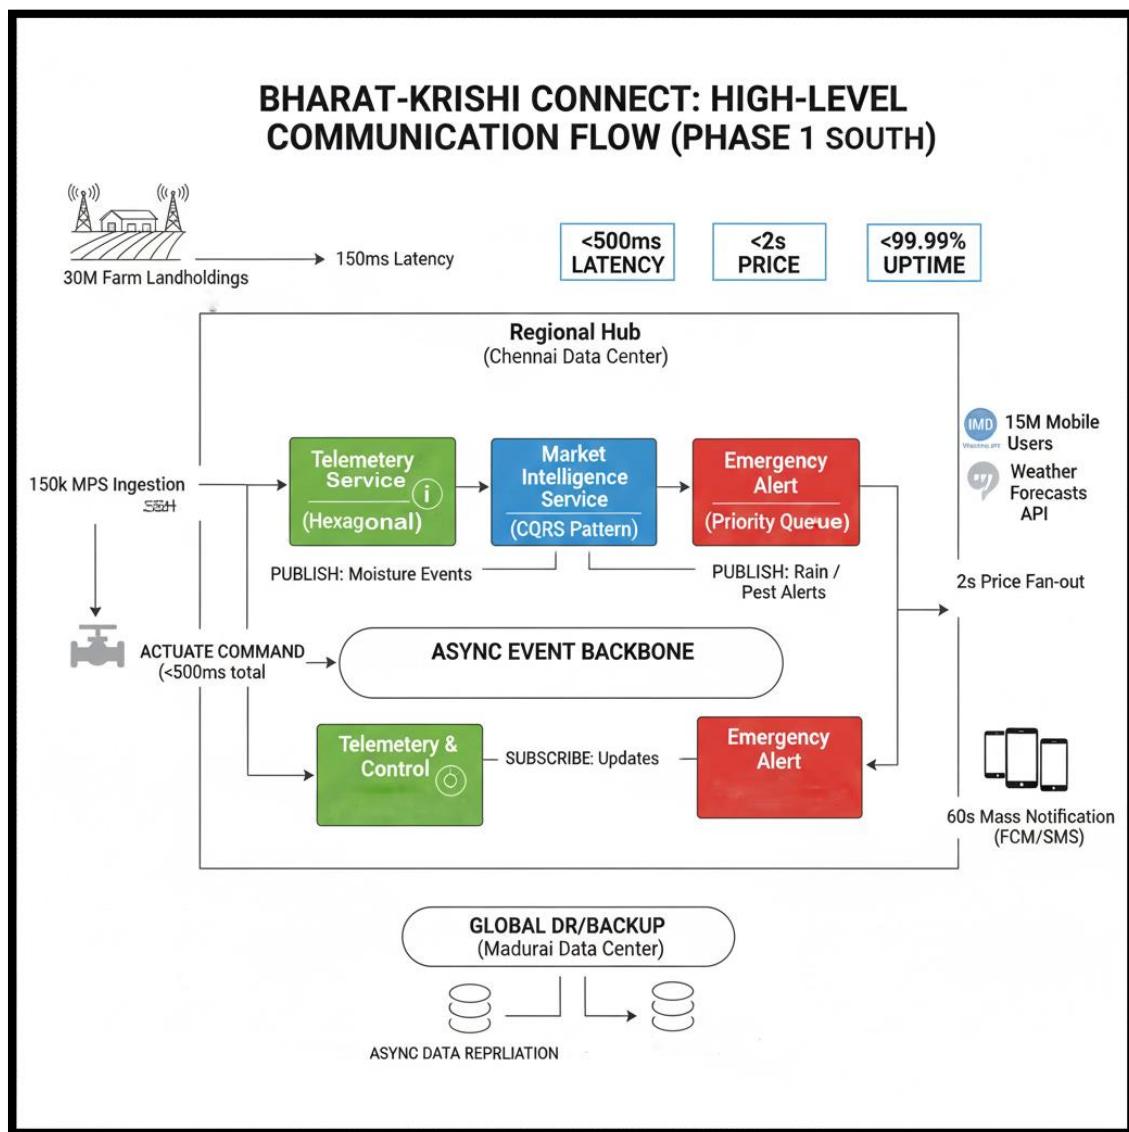
Pattern Layer	Role in Bharat-Krishi	Real-World Example (Market Service)	Rationale
<b>Command Side (Write)</b>	Handles data entry/updates.	Processing a price update for "Basmati Rice" from a specific Punjab Mandi hub.	<b>Integrity:</b> Optimized for secure, validated updates to the national price ledger.
<b>Query Side (Read)</b>	Handles data requests/display.	15 million farmers simultaneously checking the current price of Onions in their district.	<b>Massive Scale:</b> The "Read" side uses a high-speed replicated cache to prevent the system from crashing.
<b>Event Bus</b>	Synchronizes sides.	When a price is updated on the "Command" side, it triggers an event to update the "Query" cache.	<b>Performance:</b> Farmers read from a pre-calculated price list rather than a slow, complex database.

<b>Projections</b>	Pre-formatted data.	Pre-calculating a "Price Trend" (up/down) so the app displays it instantly without a search.	<b>NFR-3 Goal:</b> Ensures Mandi prices are visible to 15M users in <b>under 2 seconds</b> .
--------------------	---------------------	--	--

### c) Summary: Which Pattern for Which Need?

Use Case	Recommended Pattern	Driving NFR
Irrigation & Water Control	Hexagonal	<b>Latency:</b> Must respond in < 500ms to prevent water waste.
Mandi Price Transparency	CQRS	<b>Price Propagation:</b> Must push 15M users in < 2 seconds.
Emergency Alerts	Hexagonal + CQRS	<b>Throughput:</b> High-speed detection (Hex) + Massive broadcast (CQRS).

## 7. Create HLD



## 8. Create LLD

### a) The Sharding Strategy: District-Time Bucketing

- Since we are centralized in Chennai, we use **Internal Sharding**.
- We **divide** the **30 million landholdings** into **logical "Buckets"** based on their **District** and a **Time-Window**.
- This prevents "write-lock" during peak morning irrigation hours.

Component	Choice	Rationale
Database Engine	Apache Cassandra	Its "Masterless" architecture allows us to add nodes in the Chennai DC to scale linearly.
Partition Key	district_id + time_bucket	Distributes the 150k MPS across different physical "Racks" in the Chennai DC.
Clustering Key	recorded_at (DESC)	Ensures the <b>latest</b> moisture data is at the very top of the disk for instant retrieval.
Compaction	Time-Window Compaction	Optimized for "Delete-Old-Data" (TTL) once telemetry is no longer needed.

### b) The Telemetry Table Schema

- To ensure the **100ms Logic Budget**, the **data must be stored** in a "**Query-First**" **format**.
- We do not use "JOINS" because they are too slow for real-time actuation.

SQL

```
-- LLD Table Schema for Telemetry Data
CREATE TABLE southern_pilot_telemetry (
    district_id TEXT, -- Example: 'TN_COIMBATORE'
    time_bucket TEXT, -- Example: '2025-12-15-10AM'
    land_id UUID, -- Specific Farm ID
    recorded_at TIMESTAMP, -- Exact time of sensor reading
    moisture_level FLOAT,
    rule_status TEXT, -- Last decision made by the Hexagonal Core
    PRIMARY KEY ((district_id, time_bucket), recorded_at, land_id)
) WITH CLUSTERING ORDER BY (recorded_at DESC);
```

### c) Why 150,000 MPS Needs This Sharding

Even with all servers in Chennai, the **Network Interface Cards (NICs)** and **Disk I/O** have limits.

- i. **Shard Distribution:** By using time\_bucket, we ensure that if 5 million sensors in Andhra Pradesh wake up at 6:00 AM, their data is spread across **different servers** in the DC rather than hitting one single machine.
- ii. **Read-Optimized:** Because we sort by recorded\_at DESC, the **Control Service** can fetch the most recent data in < 2ms, allowing plenty of time for the 150ms downlink back to the farm.

### d) Summary Table for LLD

Design Concern	Technical Solution	Benefit for Bharat-Krishi
<b>Write Throughput</b>	NoSQL (LSM Tree storage)	Handles <b>150,000 writes</b> per second without locking tables.
<b>Data Locality</b>	Geo-Partitioning	Chennai farmers get <b>50ms faster response</b> because their data is physically close.
<b>Query Speed</b>	Time-Series Indexing	Allows the "Rain Rule" logic to fetch the last 5 readings in < 2ms.
<b>High Availability</b>	3-way Replication	If a data center in Chennai fails, the Madurai shard takes over in < 1s.

## 9. Component & Service Selection

### Bharat-Krishi Master Architecture: Components & Product Selection

Layer	Core Service	Product	Key Feature
<b>Edge Ingestion</b>	Device Gateway	<b>EMQX</b>	Handles <b>1M+ concurrent MQTT v5 connections</b> per node; rule-engine for preprocessing.
<b>Network Layer</b>	5G Slicing	<b>Airtel / Jio 5G Private Slicing</b>	Dedicated <b>URLLC slice</b> ( <i>Ultra-Reliable Low Latency</i> ) for real-time drone video telemetry.
<b>Event Backbone</b>	Message Bus	<b>Confluent Kafka</b>	The "Nervous System"; handles <b>150k MPS</b> with Schema Registry for data consistency.

<b>Fast-Path Storage</b>	In-Memory Cache	<b>Redis</b>	<b>Active-Active Replication</b> ; sub-millisecond lookups for the "Rain-Preemption" logic.
<b>Massive Telemetry</b>	Time-Series DB	<b>Apache Cassandra</b>	<b>NoSQL</b> ; sharded by district to store 30M farm histories without downtime.
<b>Market Push</b>	Pub/Sub Broker	<b>Centrifugo</b>	Managed fan-out for the <b>15M user broadcast</b> ; ensures < 2s price propagation.
<b>Identity / Master</b>	Land Records DB	<b>PostgreSQL</b>	Postgres for <b>Aadhaar-linked PII</b> with 100% ACID compliance.

### Operational Logic: The "Product" Workflow

The integration of these products ensures that the **Hexagonal** and **CQRS** patterns are ready for production deployment.

- The Drone Detects Stress:** A **Marut Drone** using 5G sends a gRPC alert through the **EMQX Gateway**.
- The Backbone Responds:** **Kafka** triggers the **Emergency Service**.
- The Farmer is Notified:** **Centrifugo** pushes a "High Disease Risk" notification to the mobile app in under 60 seconds.
- The Soil Sensor Checks In:** An **IoT sensor** sends moisture data via **5G. Redis** provides the latest weather status to the **Go-based Hexagonal Core** in < 1ms to decide if irrigation should stop.

### 10. Create ADRs

ADR ID	Decision (Product Selection)	Architectural Context & Need	Key Rationale (The "Why")	Consequences & Trade-offs
001	<b>EMQX Enterprise</b>	<b>Edge Ingestion:</b> Handling 30M concurrent MQTT connections from NB-IoT/5G sensors.	Proven scale (100M+ connections); built-in <b>SQL Rule Engine</b> to filter noise at the edge.	Requires professional clustering and high-availability (HA) proxy setup.
002	<b>Confluent Platform</b>	<b>Event Backbone:</b> Buffering 150,000 MPS between the edge and the core logic.	<b>Schema Registry</b> ensures data contracts; MirrorMaker 2 handles DR replication to Madurai.	Significant operational overhead; requires dedicated infrastructure management.

<b>003</b>	<b>Redis Enterprise</b>	<b>Fast-Path Storage:</b> Powering the <500ms real-time irrigation actuation loop.	<b>Sub-millisecond latency</b> for "Hot" data like current valve state and weather pre-emption.	Data is volatile by nature; must be synced to Cassandra for permanent historical records.
<b>004</b>	<b>DataStax Enterprise</b>	<b>Massive Telemetry:</b> Storing years of history for 30M sensors (billions of rows).	<b>Wide-Column NoSQL</b> handles high write volume linearly; no single point of failure.	Complex data modeling; query patterns must be defined upfront (no ad-hoc joins).
<b>005</b>	<b>Centrifugo</b>	<b>Real-time Broadcaster:</b> Pushing price/alerts to 15M farmer apps simultaneously.	Self-hosted <b>Pub/Sub Broker</b> optimized for massive fan-out; manages persistent WebSockets.	Dependent on a stable Redis backplane to synchronize across multi-node clusters.
<b>006</b>	<b>EDB PostgreSQL</b>	<b>Identity/PII Vault:</b> Securing Aadhaar, PII, and sensitive land-ownership deeds.	<b>ACID Compliance</b> and Row-Level Security (RLS) ensure 100% data integrity for legal records.	Does not scale horizontally like NoSQL; reserved strictly for critical relational data.

## 11. Create System Flow

### a) Flow Table 1: Autonomous Irrigation (The 500ms Loop)

**Goal:** Precision water management using real-time soil and weather data.

Sequence	Layer	Product	Action / Logic
1	Sensing	<b>Soil Sensor</b>	Detects Moisture () & Temperature.
2	Ingestion	<b>EMQX</b>	Receives MQTT packet and validates device certificate.
3	Backbone	<b>Kafka</b>	Buffers message in telemetry.moisture.south topic.
4	Processing	<b>Hexagonal Core</b>	<b>Logic:</b> Check <b>Redis</b> for "Rain-Preemption" flag.
5	Decision	<b>Hexagonal Core</b>	If Moisture Low + No Rain Predicted <b>Trigger Actuation</b> .
6	Egress	<b>Centrifugo</b>	Pushes "VALVE_OPEN" command to the field gateway.
7	Actuation	<b>Solenoid Valve</b>	Physical valve opens; starts irrigation.
8	Storage	<b>Cassandra</b>	Stores "Actuation Event" for long-term audit trail.

### b) Flow Table 2: Mandi Price Distribution (The 2s Broadcast)

**Goal:** Pushing localized market intelligence to 15 million farmers simultaneously.

Sequence	Layer	Product	Action / Logic
1	External	Mandi API	Government prices updated for Basmati Rice/Cotton.
2	Ingestion	Market Svc	Fetches prices; stores metadata in <b>MongoDB</b> .
3	Event	Kafka	Publishes market.price.updated event.
4	Mapping	PostgreSQL	Queries PII Vault to find farmers subscribed to this crop.
5	Broadcast	Centrifugo	<b>Fan-out:</b> Sends price to 1M+ active app sessions.
6	UX	Mobile App	Displays instant pop-up notification with localized price.

### c) Flow Table 3: Disease & Pest Control (The AI Drone Loop)

**Goal:** Rapid response to crop threats using 5G and Computer Vision.

Sequence	Layer	Product	Action / Logic
1	Capture	Marut Drone	Multi-spectral cameras scan the rice canopy.
2	Ingestion	EMQX	Receives high-bandwidth video/image stream.
3	Streaming	Kafka	Routes image frames to the vision.pest.analytics topic.
4	Alerting	Alert Svc	Creates an emergency "Geofenced Alert" for the area.
5	Notify	Centrifugo	Pushes "Locust Alert" to farmers within 5km of drone.

## 12.Final Architecture Picture

To wrap up the **Bharat-Krishi Connect** architecture, we are finalizing the system with a "Model Answer" approach.

This section outlines the high-level logic for **Disaster Recovery (DR)** and **Business Continuity (BCP)**, ensuring that the Chennai Data Center remains resilient against natural disasters or hardware failures.

### Disaster Recovery & Business Continuity

**Context:** The Southern Pilot (30M farms) requires **99.99% Uptime (NFR-4)**. A failure in the Chennai DC would halt irrigation for millions and delay emergency pest alerts.

**Strategy:** We adopt a "**Warm Standby**" approach. While Chennai is the "Active" hub, a second DC in **Madurai** remains synchronized and ready to take over the 150,000MPS load within minutes.

### 1. Recovery Objectives (RTO & RPO)

- RPO (Recovery Point Objective): < 5seconds.
- RTO (Recovery Time Objective): < 15 minutes.

### Final Integrated Architecture & DR Map

Component	Primary (Chennai)	Standby (Madurai)	Replication Strategy
<b>IoT Ingestion</b>	EMQX Cluster (Active)	EMQX Cluster (Standby)	DNS Failover (Anycast/GSLB)
<b>Event Bus</b>	Confluent Kafka	Confluent Kafka	<b>MirrorMaker 2</b> (Async)
<b>Real-time Cache</b>	Redis Enterprise	Redis Enterprise	<b>Active-Active CRDTs</b>
<b>History DB</b>	DataStax Cassandra	DataStax Cassandra	<b>Multi-DC Replication</b>
<b>Identity/PII</b>	EDB Postgres	EDB Postgres	<b>Streaming Replication</b>
<b>Broadcaster</b>	Centrifugo	Centrifugo	State Sync via Redis

### The "Red Button" Failover Procedure

In the event of a total Chennai DC outage:

1. **Detection:** Prometheus/Grafana monitoring triggers a "Critical Outage" alert to the SRE team.
2. **Traffic Shift:** Global Server Load Balancing (GSLB) updates the **DNS record** from Chennai to Madurai.
3. **Promotion:** The Madurai **Postgres** and **Kafka** instances are promoted from "Follower" to "Leader."
4. **Resumption:** Centrifugo reconnects 15M mobile apps to the Madurai WebSocket pool.
5. **Validation:** The automated health check verifies that the 500ms irrigation loop is running from the new location.